

Midterm and automatic differentiation (aka autograd)

Lecture 06 — CS 577 Deep Learning

Instructor: Yutong Wang

Computer Science
Illinois Institute of Technology

September 25, 2024

Topics

- The midterm
- Attention models
- Tricks for doing homework 3.
 - `np.matmul` with multidim arrays
 - Batched matrix multiplication
 - Batched diagonalization
 - Batched outer product
 - `np.moveaxis`
- Automatic differentiation

Midterm ↑
exam ↓

More on homework 3: Next-word-prediction

φ word embedding $\varphi(\text{over})$ $\varphi(\text{the}) \in \mathbb{R}^d$

The quick brown fox jumps over the -----

$\mathbf{X}^{(i)} = [\mathbf{x}^{(i,1)} \quad \mathbf{x}^{(i,2)} \quad \dots \quad \mathbf{x}^{(i,C-1)} \quad \mathbf{x}^{(i,C)}] \in \mathbb{R}^{d \times C}$ is a $d \times C$ matrix

↑
a single data point (also called a prompt)

$y^{(i)} \in \{1, \dots, K\}$ is set of all candidate words
the
vocabulary

More on homework 3: Next-word-prediction

The quick brown fox jumps over the ____.

$\mathbf{X}^{(i)} = [\mathbf{x}^{(i,1)} \quad \mathbf{x}^{(i,2)} \quad \dots \quad \mathbf{x}^{(i,C-1)} \quad \mathbf{x}^{(i,C)}] \in \mathbb{R}^{d \times C}$ is a $d \times C$ matrix

empty

the
should have the most important

$y^{(i)} \in \{1, \dots, K\}$ is set of all candidate words

More on homework 3: Next-word-prediction

The quick brown fox jumps over the -----.

$\mathbf{X}^{(i)} = [\mathbf{x}^{(i,1)} \quad \mathbf{x}^{(i,2)} \quad \dots \quad \mathbf{x}^{(i,C-1)} \quad \mathbf{x}^{(i,C)}] \in \mathbb{R}^{d \times C}$ is a $d \times C$ matrix

- unrealistic

- good for learning + writing paper

$y^{(i)} \in \{\pm 1\}$ we are keeping it simple

Notations

- C as `n_context` — dim of word embed 700-ish
- d as `n_features` — training
- n as `n_samples` —
- q as `n_reduced`, where $q < d$ — projection onto small space of dim q

Let

$$\theta = [W^{(1)}, W^{(2)}, w^{(3)}]$$

where


$$W^{(1)} \text{ and } W^{(2)} \in \mathbb{R}^{q \times d}$$

and

$$w^{(3)} \in \mathbb{R}^d$$

and

$$f(X^{(i)}; \theta) = w^{(3)\top} X^{(i)} \underbrace{\text{softmax} \left(X^{(i)\top} W^{(2)\top} W^{(1)} x^{(i,C)} \right)}_{\text{attention}}.$$




subspace
of some
concept.

(“Single-layer” and “single-headed”) Attention

$$f(X^{(i)}; \theta) = w^{(3)\top} X^{(i)} \text{softmax} \left(X^{(i)\top} W^{(2)\top} W^{(1)} x^{(i,C)} \right).$$

\uparrow
 $[x^{(i,1)} \dots x^{(i,C)}]$




\sum to 1 and positive

$\propto \exp(x^{(i,j)})^\top W^{(2)\top} W^{(1)} x^{(i,C)}$
 prop to
 Similarity between
 j-th word emb
 &
 c-th word emb


! Note softmax playing totally different role compared to CE.

("Single-layer" and "single-headed") Attention

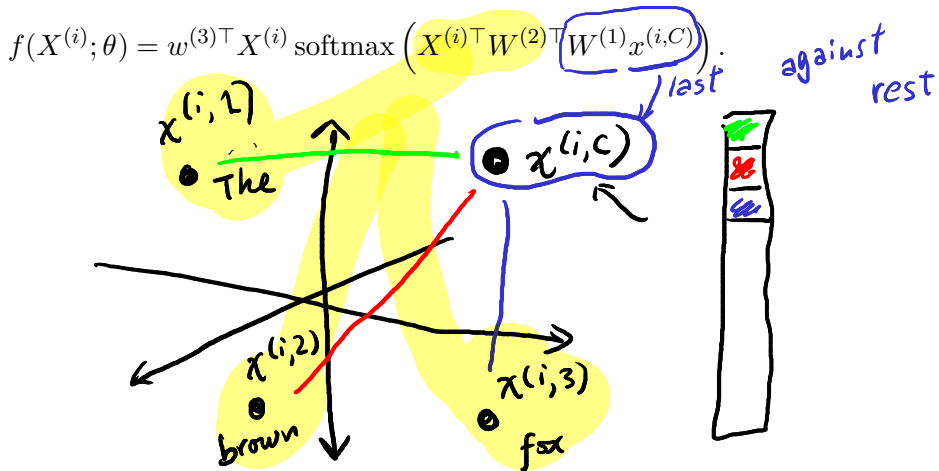
$$f(X^{(i)}; \theta) = w^{(3)\top} X^{(i)} \text{softmax} \left(X^{(i)\top} W^{(2)\top} W^{(1)} x^{(i,C)} \right).$$




(“Single-layer” and “single-headed”) Attention



(“Single-layer” and “single-headed”) Attention




(“Single-layer” and “single-headed”) Attention

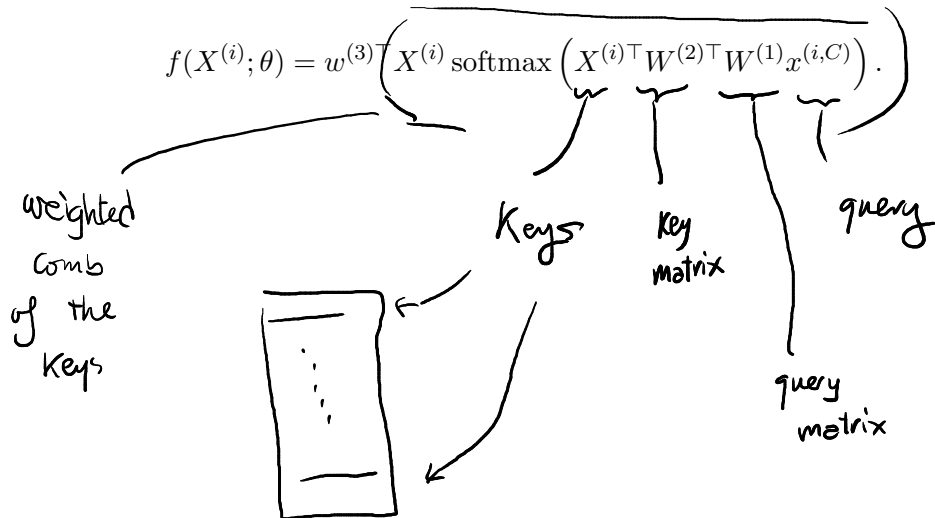


(“Single-layer” and “single-headed”) Attention

$$f(X^{(i)}; \theta) = w^{(3)\top} X^{(i)} \text{softmax} \left(X^{(i)\top} W^{(2)\top} W^{(1)} x^{(i,C)} \right).$$




(“Single-layer” and “single-headed”) Attention




(“Single-layer” and “single-headed”) Attention

$$f(X^{(i)}; \theta) = w^{(3)\top} X^{(i)} \text{softmax} \left(X^{(i)\top} W^{(2)\top} W^{(1)} x^{(i,C)} \right).$$



(“Single-layer” and “single-headed”) Attention

$$f(X^{(i)}; \theta) = \underline{w}^{(3)\top} X^{(i)} \text{softmax} \left(X^{(i)\top} W^{(2)\top} W^{(1)} x^{(i,C)} \right).$$



Reason for "(1), (2)" is to reuse
W for weights

(“Single-layer” and “single-headed”) Attention

$$f(X^{(i)}; \theta) = w^{(3)\top} X^{(i)} \text{softmax} \left(X^{(i)\top} W^{(2)\top} W^{(1)} x^{(i,C)} \right). \quad \leftarrow \text{typo}$$

bold W capital stands multi-dim arr
including matrix & tensor

bold w lower case = vector

non-bold w l.c. = scalar

In this class
Goodfellow
convention

(“Single-layer” and “single-headed”) Attention

$$f(X^{(i)}; \theta) = w^{(3)\top} X^{(i)} \underset{\text{sign}}{\text{softmax}} \left(X^{(i)\top} W^{(2)\top} W^{(1)} x^{(i,C)} \right).$$

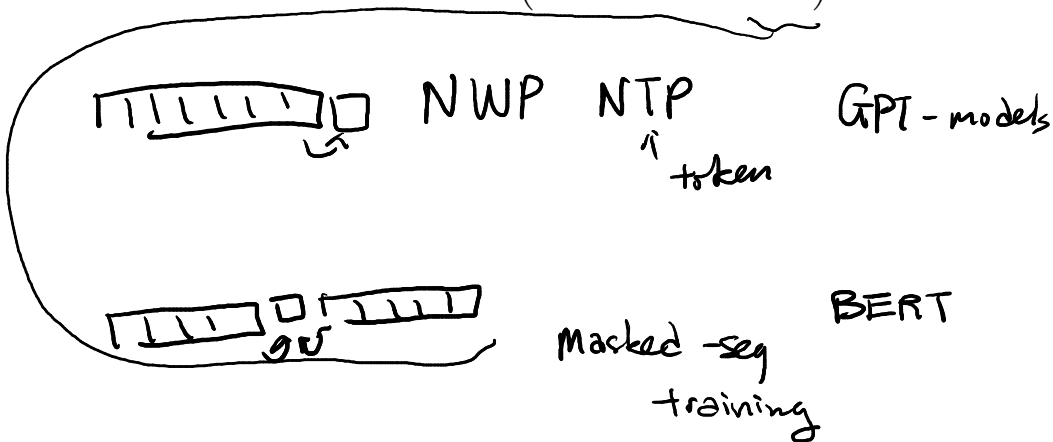
C - dot products

$\{ \pm 1 \}$ $+1 = \text{foo}$
 $-1 = \text{bar}$

$\{ 1, \dots, K \}$ $\hookrightarrow \text{argmax}$ $1 = \text{the}$
 $2 = \text{an}$
 $3 = \text{and}$

(“Single-layer” and “single-headed”) Attention

$$f(X^{(i)}; \theta) = w^{(3)\top} X^{(i)} \text{softmax} \left(X^{(i)\top} W^{(2)\top} W^{(1)} x^{(i,C)} \right).$$



Tricks for doing homework 3.

- `np.matmul` with multidim arrays
- Batched matrix multiplication
- Batched diagonalization
- Batched outer product
- `np.moveaxis`

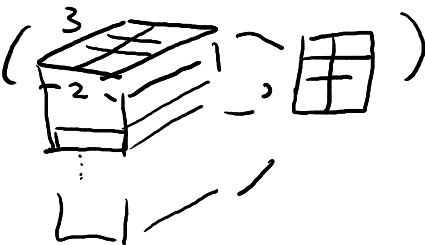
np.matmul with multidim arrays

~~X~~ rank 3
tensor

$W^{(2)}$ rank 2
tensor

```
1 # IN (n_samples, n_context, n_features) @ (n_features, n_reduced)
2 Keys = np.matmul(X, W2)
3 # OUT (n_samples, n_context, n_reduced)
```

aka
matrix

np.matmul ()



a stack of
matmul

Batched ~~matrix~~ multiplication

Keys

queries.shape

```
1 queries.shape == (n_samples, n_reduced)
2
3 # (n_samples, n_context, n_reduced)
4 np.sum(
5     # (n_samples, n_context, n_reduced) * (n_samples, 1, n_reduced)
6     Keys * np.expand_dims(queries, axis=1),
7     # --> (n_samples, n_context, n_reduced)
8     axis=2)
9 # --> (n_samples, n_context)
```

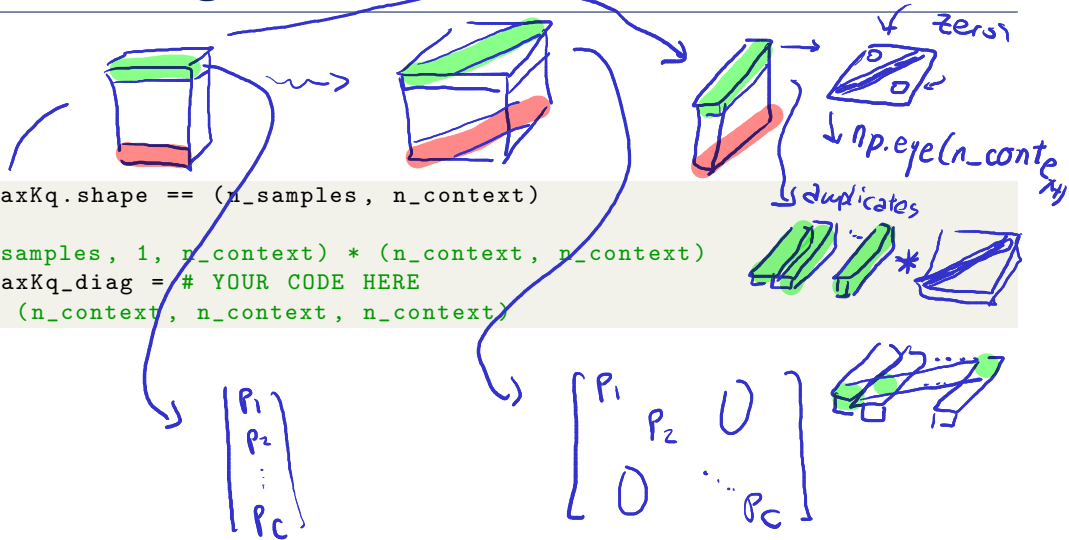


match the
0th dim
and do
stackwise matmul

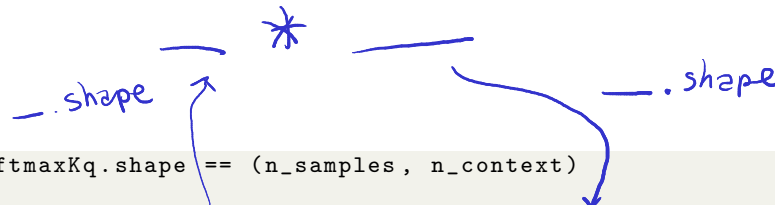
sum(axis=2)

Batched diagonalization

```
1 softmaxKq.shape == (n_samples, n_context)
2
3 # (n_samples, 1, n_context) * (n_context, n_context)
4 softmaxKq_diag = # YOUR CODE HERE
5 # --> (n_context, n_context, n_context)
```




Batched diagonalization




```
1 softmaxKq.shape == (n_samples, n_context)
2
3 # (n_samples, 1, n_context) * (n_context, n_context)
4 softmaxKq_diag = # YOUR CODE HERE
5 # --> (n_context, n_context, n_context)
```


Batched outer product



```
1 # (n_samples, 1, n_context) * (n_samples, n_context, 1)
2 softmaxKq_outer = # YOUR CODE HERE
3 # --> (n_samples, n_context, n_context)
```



Batched outer product




`expand_dims(axis=1)`

```
1 # (n_samples, 1, n_context) * (n_samples, n_context, 1)
2 softmaxKq_outer = # YOUR CODE HERE
3 # --> (n_samples, n_context, n_context)
```

`n_context`

`n_context`




np.moveaxis

$X^{(i)}$ stacked

not touching D

```
1 X.shape == (n_samples, n_context, n_features)
2
3 # (n_samples, n_features, n_context) @ (n_samples, n_context, n_context)
4 A = np.moveaxis(X, 1, 2) @ D
5 # --> (n_samples, n_features, n_context)
```



want batchwise matmul

transpose each stack

Welcome to the second half the course

- Understand how automatic differentiation works under the hood
- What is overfitting, when it happens, and what can you do about it
- Transfer learning
- Convolution layers, attention
- How to compute things in deep learning really really fast.

Automatic differentiation


Let's compute

$$f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$$


$$x, y \in \mathbb{R}$$

using automatic differentiation.


Eval $f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ at $(x, y) = (2, 3)$




Eval $f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ **at** $(x, y) = (2, 3)$




Eval $f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ **at** $(x, y) = (2, 3)$




Eval $f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ **at** $(x, y) = (2, 3)$



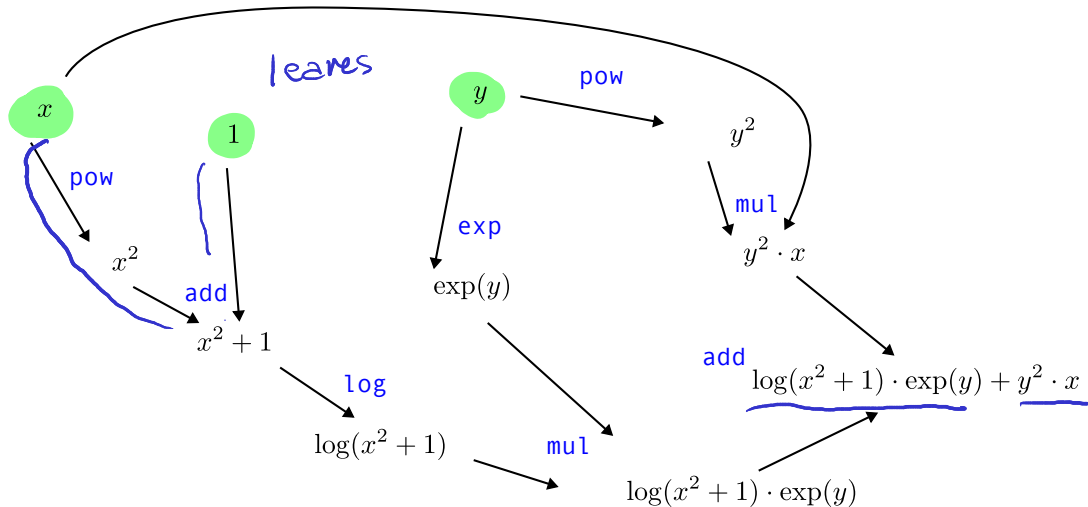
Eval $f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ **at** $(x, y) = \underline{(2, 3)}$




Eval $f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ **at** $(x, y) = (2, 3)$




Eval $\nabla f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ **at** $(2, 3)$




Eval $\nabla f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ at $(2, 3)$




Eval $\nabla f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ at $(2, 3)$



Eval $\nabla f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ at $(2, 3)$




Eval $\nabla f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ **at** $(2, 3)$



An autograd “module”

```
1 class ag: # AutoGrad
2     class Scalar: # Scalars with grads
3         def __init__(self, value, op="", _backward= lambda : None,
4             inputs=[], label=""):
5             self.value = float(value)
6             self.grad = 0.0
7
8             # ... lines skipped
```




An autograd “module”

```
1 class ag: # AutoGrad
2     class Scalar: # Scalars with grads
3         def __init__(self, value, op="", _backward= lambda : None,
4             inputs=[], label=""):
5             self.value = float(value)
6             self.grad = 0.0
7
8             # ... lines skipped
9
10        def __add__(self, other): # ...
11        def __mul__(self, other): # ...
12        def __pow__(self, exponent): # ...
13
14    def exp(input): # ...
15    def log(input): # ...
```


Implement f in our framework

```
1 x = ag.Scalar(2, label="z2:x") # wrap x
2 y = ag.Scalar(3, label="z3:y") # wrap y
3
4 # implement log(x^2+1)*exp(y) + y^2*x
5 def f(x,y):
6     z1 = ag.Scalar(1, label="z1:1") # constant 1
7     z2 = x
8     z3 = y
9
10    z4 = z2**2
11    z4.label = "z4:x^2"
12
13    # ... lines skipped
```


Forward pass — DFS — computation graph



Forward pass — DFS — computation graph



Forward pass — DFS — computation graph



Exercise 2.A

Implement `ag.Scalar.topological_sort`


```
1 def topological_sort(self):  
2     ## EXERCISE 2.A  
3     ## YOUR CODE HERE
```

```
1 Check your answer  
2 [z2:x,  
3  z4:x^2      pow(2),  
4  z1:1,  
5  z5:x^2 + 1   add,  
6  z6:log(x^2 + 1) log,  
7  z3:y,  
8  z7:exp(y)     exp,  
9  z8:log(x^2 + 1) * exp(y) mul,  
10 z9:y^2      pow(2),  
11 z10:y^2 * x   mul,  
12 z11:log(x^2 + 1) * exp(y) + y^2 * x add]
```

← leaves first .

← root last

Eval $\nabla f(x, y) = \log(x^2 + 1) \exp(y) + y^2 x$ at $(2, 3)$



Backward and add

```
1 def backward(self):
2     self.grad = 1.0
3     topo_order = self.topological_sort()
4     for node in reversed(topo_order):
5         node._backward()
6
7 def __add__(self, other):
8     assert isinstance(other, ag.Scalar)
9     output = ag.Scalar(self.value + other.value,
10                        inputs=[self, other], op="add")
11     def _backward():
12         # pass
13         self.grad += output.grad 1.0
14         other.grad += output.grad 1.0
15
16     output._backward = _backward
17     return output
```


Exercise 2.B

Implement `ag.Scalar.topological_sort`

```
1  def __mul__(self, other):
2      # ...
3      def _backward():
4          ## EXERCISE 2.B
5          ## YOUR CODE HERE
6          return None
7  # ... and also the backward function in
8      def __pow__(self, exponent): # exponent is just a python float
9  # ... and in
10     def exp(input):
```

*self * other*

*x ** exponent*

```
1  Check your answer
2  f(x,y) 50.3264246157732
3  x.grad 25.068429538550134
4  y.grad 44.3264246157732
```

Exercise 2.C

Implement `ag.Scalar.relu`

```
1 def relu(input):
2     output = ag.Scalar(max(0, input.value), inputs=[input], op="relu")
3
4     def _backward():
5         ## EXERCISE 2.C
6         ## YOUR CODE HERE
7
8 # ...
9
10 # Check your answer with the plot
11 xs_raw = [i - 5 for i in range(10)]
12 xs = [ag.Scalar(i) for i in xs_raw]
13 ys = [ag.relu(x)**3 for x in xs]
14 [y.backward() for y in ys]
15 grads = [x.grad for x in xs]
16 plt.plot(xs_raw, grads, label='autograd')
17 plt.plot(xs_raw, [3*max(0,x_raw)**2 for x_raw in xs_raw], label='manual grad')
```

What can we do with `ag.Scalar`?

For the remaining of this class, we will build a “knock-off” PyTorch consisting of:

- A `Model` class that encapsulates the parameters and the forward pass
- A `Loss` class for calculating Mean Squared Error (MSE)
- An `Optimizer` class for performing gradient descent
- A training loop

We will fit a 1-hidden layer neural network

$$f(x; \mathbf{w}_1, \mathbf{b}_1, \mathbf{w}_2, b_2) = \mathbf{w}_2^\top \text{relu}(\mathbf{w}_1 x + \mathbf{b}_1) + b_2$$

where \mathbf{w}_1 and \mathbf{w}_2 are both vectors of `n_hidden` dimensions.

The Model class

```
1 class Model:
2     def __init__(self, n_hidden, rng_seed=42):
3         np.random.seed(rng_seed)
4
5         w1np = np.random.randn(n_hidden)
6         # ...
7         b2np = np.random.randn(1)
8
9         self.w1 = [ag.Scalar(val) for val in w1np]
10        # ...
11        self.b2 = [ag.Scalar(val) for val in b2np]
12
13        self.parameters = self.w1 + self.b1 + self.w2 + self.b2
14
15    def forward(self, x): # ...
```

The forward function

```
1  def forward(self, x):
2      # "upgrade" x into ag.Scalars
3      x_scalar = [ag.Scalar(val) for val in x]
4      n_samples = len(x_scalar)
5
6      # calculate the forward
7
8      ## YOUR CODE HERE
9      return [ag.Scalar(0.0) for i in range(n_samples)]
```

The Loss class

```
1 class Loss:
2     def mse(self, predictions, targets):
3         # mean squared error
4         assert len(predictions) == len(targets)
5         n_samples = len(predictions)
6         loss = ag.Scalar(0.0)
7
8         # YOUR CODE HERE
9
10    return loss
```


The Optimizer class

```
1 class Optimizer:
2     def __init__(self, parameters, lr=0.01):
3         self.parameters = parameters
4         self.lr = lr
5
6     def zero_grad(self):
7         # YOUR CODE HERE
8         pass
9
10    def step(self):
11        # YOUR CODE HERE
12        pass
```

Training Loop

```
1 model = Model(n_hidden=20)
2 loss_fn = Loss()
3 optimizer = Optimizer(model.parameters, lr=0.1)
4
5 for epoch in range(100):
6     optimizer.zero_grad()
7     output = model.forward(xnp)
8     loss = loss_fn.mse(output, ynp)
9     loss.backward()
10    optimizer.step()
11    if epoch % 10 == 0:
12        print(f"Iteration {epoch}, Loss: {loss.value}")
```


Exercise 3



References I
