

Convolutional nets

Lecture 08 — CS 577 Deep Learning

Instructor: Yutong Wang

Computer Science
Illinois Institute of Technology

October 9, 2024

Autograd-enabled tensors

Previously

- `ag.Scalar`

Today:

- Use `ag.Scalar` for training small relu networks
- `ag.Tensor`
- ~~Overflow in cross entropy and how to avoid it~~
- Use `ag.Tensor` to train convolutional neural networks (CNNs)

Autograd-enabled scalars: `ag.Scalar`

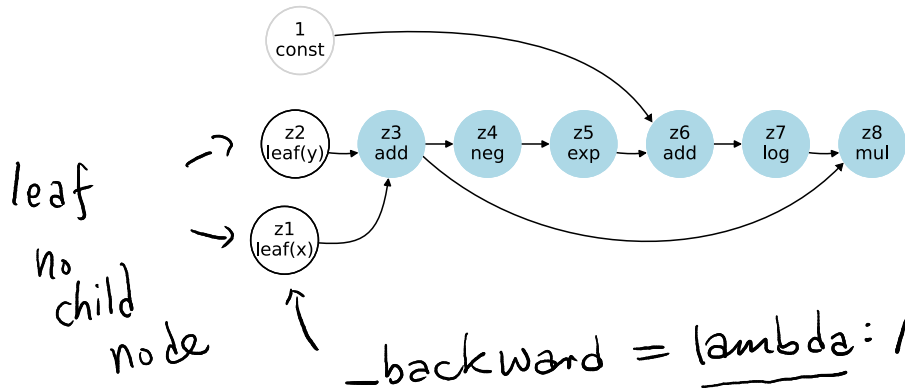
```
1  class Scalar: # Scalars with grads
2      def __init__(self,
3                    value,
4                    op="",
5                    _backward= lambda : None,
6                    inputs=[],
7                    label=""):
8
9      self.value = float(value)
10     self.grad = 0.0
11
12     self._backward = _backward
13     self.inputs = inputs
14
15     self.op = op
16     self.label = label
```

upgraded
number
with info
about
the grad

Computational graph

$$f(x, y) = \log(1 + e^{-(x+y)}) \cdot (x + y).$$

restricted to be
scalar



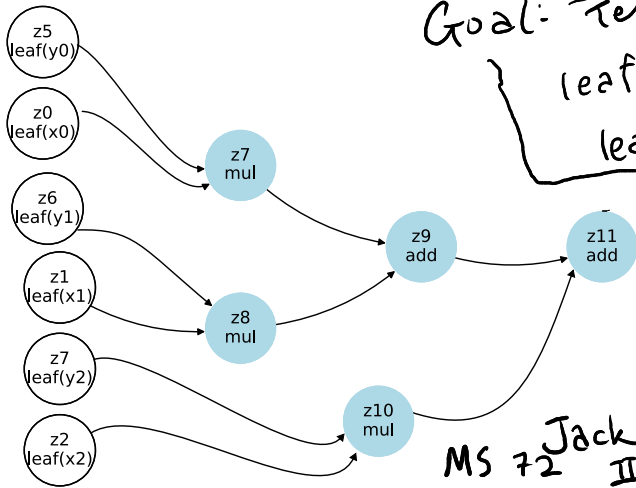
Computational graph

$$f(\mathbf{x}, \mathbf{y}) = \mathbf{x}^\top \mathbf{y} = x_1 y_1 + x_2 y_2 + x_3 y_3$$

brute force calculate
of dot product

Goal: Tensor quant
leaf(x) → Matmul
leaf(y) →

Cuda's
BLAS




can leverage
CPU → BLAS
Basic
Linear
Algebra
Subr

MS 72 Jack Dongarra
IIT

Computational graph

$$-\log(e^{z_1}/(e^{z_1} + e^{z_2} + e^{z_3})) = L_{CE}\left(\begin{bmatrix} z_1 \\ z_2 \\ z_3 \end{bmatrix}, 1\right)$$



classify
with 3 category

NLL

```
1 expz1 = ag.exp(z1)
2 expz2 = ag.exp(z2)
3 expz3 = ag.exp(z3)
4 final_output = -ag.log( expz1/(expz1+expz2+expz3))
```

What can we do with `ag.Scalar`?

PyTorch-“lite” consisting of:

- A `Model` class (params and forward)
- A `Loss` class (just Mean Squared Error (MSE) for now)
- An `Optimizer` class (just gradient descent for now)
- A training loop

We will fit a 1-hidden layer neural network with 1-dimensional input

$$f(x; \mathbf{w}_1, \mathbf{b}_1, \mathbf{w}_2, b_2) = \mathbf{w}_2^\top \text{relu}(\mathbf{w}_1 x + \mathbf{b}_1) + b_2$$

where \mathbf{w}_1 and \mathbf{w}_2 are both vectors of `n_hidden` dimensions.

Training Loop

```
1 model = Model(n_hidden=20)
2 loss_fn = Loss()
3 optimizer = Optimizer(model.parameters, lr=0.1)
4
5 for epoch in range(100):
6     optimizer.zero_grad()
7     output = model.forward(xnp)
8     loss = loss_fn.mse(output, ynp)
9     loss.backward()
10    optimizer.step()
11    if epoch % 10 == 0:
12        print(f"Iteration {epoch}, Loss: {loss.value}")
```

*Pytorch / JAX /
TensorFlow*

Exercise 1

[lec08-in-class-ex1-framework.ipynb](#)


$w_1 \in \mathbb{R}^{10}$



$x^{(1)}, \dots, x^{(N)} \in \mathbb{R}$

$$f(x; \mathbf{w}_1, \mathbf{b}_1, \mathbf{w}_2, b_2) = \mathbf{w}_2^\top \text{relu}(\mathbf{w}_1 x + \mathbf{b}_1) + b_2$$

where \mathbf{w}_1 and \mathbf{w}_2 are both vectors of `n_hidden` dimensions.



$$\left(y^{(i)} - f(x^{(i)}; \theta) \right)^2$$

The Model class

```
1 class Model:
2     def __init__(self, n_hidden, rng_seed=42):
3         np.random.seed(rng_seed)
4
5         w1np = np.random.randn(n_hidden)
6         # ...
7         b2np = np.random.randn(1)
8
9         self.w1 = [ag.Scalar(val) for val in w1np]
10        # ...
11        self.b2 = [ag.Scalar(val) for val in b2np]
12
13        self.parameters = self.w1 + self.b1 + self.w2 + self.b2
14
15    def forward(self, x): # ...
16        # x is a 1-dimensional numpy array
```

} numpy arrays

} list of ag.Scalar

The forward function

```
1  def forward(self, x):
2      # x is a 1-dimensional numpy array
3      # "upgrade" x into ag.Scalars
4      x_scalar = [ag.Scalar(val) for val in x]
5      n_samples = len(x_scalar)
6
7      # calculate the forward
8
9      ## YOUR CODE HERE
10     return [ag.Scalar(0.0) for i in range(n_samples)]
```

↑
Task 1: change to the correct output

The Loss class

```
1 class Loss:
2     def mse(self, predictions, targets):
3         # mean squared error
4         assert len(predictions) == len(targets)
5         n_samples = len(predictions)
6         loss = ag.Scalar(0.0)
7
8         # YOUR CODE HERE
9
10    return loss
```

The Optimizer class

```
1 class Optimizer:
2     def __init__(self, parameters, lr=0.01):
3         self.parameters = parameters
4         self.lr = lr
5
6     def zero_grad(self):
7         for param in self.parameters:
8             param.grad = 0.0
9
10    def step(self):
11        for param in self.parameters:
12            param.value -= self.lr * param.grad
```

Check gradients!

```
1 Gradients w.r.t...
2 {'w1': array([-0.77812733, -0.          , ...
3
4  'b1': array([ 2.35879461, -0.          , ...
5
6  'w2': array([ 4.15815372e+00, -0.00000000e+00, ...
7
8  'b2': 3.1941792242464686}
```

Upshot

- It seems to work just fine.

Autograd-enabled tensors

Previously

- `ag.Scalar`

Today:

- ☒ Use `ag.Scalar` for training small relu networks
- `ag.Tensor`
- Use `ag.Tensor` to train convolutional neural networks (CNNs)
- Overflow in cross entropy and how to avoid it

ag.Scalar

```
1  class Scalar: # Scalars with grads
2      def __init__(self,
3                    value,
4                    op="",
5                    _backward= lambda : None,
6                    inputs=[],
7                    label=""):
8
9      self.value = float(value)
10     self.grad = 0.0
```


ag.Tensor

```
1 class Tensor: # Tensor with grads
2     def __init__(self,
3                   value,
4                   op="",
5                   _backward= lambda : None,
6                   inputs=[],
7                   label=""):
8
9         if type(value) in [float ,int]:
10             value = np.array(value)
11         self.value = value # <-----\
12 #                                     |- same shape!
13         self.grad = np.zeros_like(value) # <-----/
```

★ $\text{self.value.shape} == \text{self.grad.shape}$

Computational graph

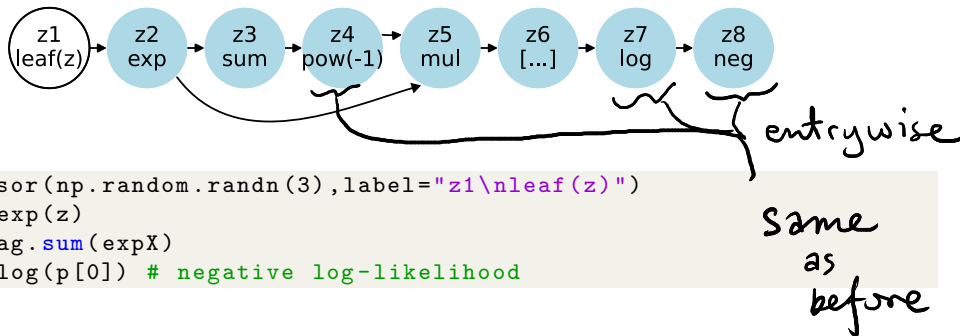
$$-\log(e^{z_1}/(e^{z_1} + e^{z_2} + e^{z_3}))$$



```
1 expz1 = ag.exp(z1)
2 expz2 = ag.exp(z2)
3 expz3 = ag.exp(z3)
4 final_output = -ag.log( expz1/(expz1+expz2+expz3))
```

Computational graph


$$-\log(e^{z_1}/(e^{z_1} + e^{z_2} + e^{z_3})) = L(\mathbf{z}, 1) \quad \leftarrow \text{cross entropy}$$



```
1 z = ag.Tensor(np.random.randn(3), label="z1\nleaf(z)")
2 expX = ag.exp(z)
3 p = expX / ag.sum(expX)
4 nll = -ag.log(p[0]) # negative log-likelihood
```

Computational graph

$$-\log(e^{z_1}/(e^{z_1} + e^{z_2} + e^{z_3})) = L(\mathbf{z}, 1) \quad \leftarrow \text{cross entropy}$$



```
1 z = ag.Tensor(np.random.randn(3), label="z1\nleaf(z)")
2 expX = ag.exp(z)
3 p = expX / ag.sum(expX)
4 nll = -ag.log(p[0]) # negative log-likelihood
```


challenge:
we broadcast

Not matmul
but
entrywise

unique to
working tensor
w₁

Computational graph

$$-\log(e^{z_1}/(e^{z_1} + e^{z_2} + e^{z_3})) = L(\mathbf{z}, 1) \quad \leftarrow \text{cross entropy}$$



```

1 z = ag.Tensor(np.random.randn(3), label="z1\nleaf(z)")
2 expX = ag.exp(z)
3 p = expX / ag.sum(expX)
4 nll = -ag.log(p[0]) # negative log-likelihood
  
```

$f(x^{(i)}; \theta)$

$z = \begin{bmatrix} z_1 \\ z_2 \\ z_3 \end{bmatrix}$



$\exp(z)$

\rightarrow


$\exp(z[1])$

$\frac{\exp(z[y[i]])}{\sum(\exp(z))}$

Computational graph

Entrywise ops

ex2.






```
1 def log(input):  
2     output = ag.Tensor(np.log(input.value), inputs=[input], op="log")  
3     def _backward():  
4         input.grad += output.grad / input.value  
5         return None  
6     output._backward = _backward  
7     return output
```

np.log
entrywise div
entrywise op

Computational graph

Slicing



```

1 def __getitem__(self, idx):
2     output = ag.Tensor(np.array(self.value[idx]),
3                           inputs = [self],
4                           op=f"[...]")

```

```


5     def _backward():
6         self.grad[idx] += output.grad # idx must not have repeats!
7         return None
8     output._backward = _backward
9     return output

```


output gradient
only flows to
input on the slice

Computational graph

Slicing



[duplicated indices]




```
1 def __getitem__(self, idx):
2     output = ag.Tensor(np.array(self.value[idx]),
3                           inputs = [self],
4                           op=f"[...]")
5
6     def _backward():
7         self.grad[idx] += output.grad # idx must not have repeats!
8         return None
9     output._backward = _backward
10    return output
```

↑ okay with dupe

Not okay.

Computational graph

Pairwise ops — multiplication (NOT matmul)



```
1 def __mul__(self, other):
2     # ... lines skipped
3     output = ag.Tensor(self.value * other.value,
4                          inputs=[self, other], op="mul")
5     def _backward():
6         self.grad += unbroadcast(output.grad*other.value, ax1, pad1)
7         other.grad += unbroadcast(output.grad*self.value, ax2, pad2)
```

broadcast
↓
built

obstacle to copying
our scalar code
is
broadcasting


What is “unbroadcast” and why define backward this way? ... next

Computational graph

Pairwise ops — multiplication (NOT matmul)

What is “unbroadcast” and why define backward this way?

- `X1.shape = (2,3,1,5,6)`
- `X2.shape = (4,1,6)`
- What is `(X1 * X2).shape`?



Need to contract


output.grad to be compatible with `X1`

Computational graph

Pairwise ops — multiplication (NOT matmul)


What is “unbroadcast” and why define backward this way?

- $X1.shape = (2, 3, 1, 5, 6)$
- $X2.shape = (4, 1, 6)$
- What is $(X1 * X2).shape$?



Computational graph


Reductive ops — sum



```
1 def sum(input,axis = None, keepdims = False):
2     output = ag.Tensor(np.sum(input.value, axis = axis, keepdims =
3     # ...
4     # ...
5     # ...
6     # ...
7     # ...
8     # ...
9     # ...
10    # ...
```

Computational graph

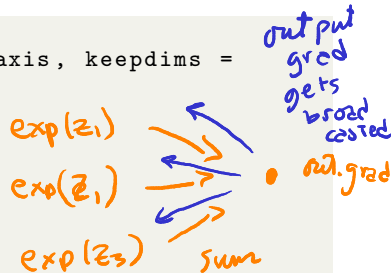
Reductive ops — sum



```
1 def sum(input,axis = None, keepdims = False):
2     output = ag.Tensor(np.sum(input.value, axis = axis, keepdims =
3     keepdims), inputs = [input], op='sum')
4     def _backward():
5         if axis == None or keepdims:
6             input.grad += output.grad
7         else:
8             # ...
9             # ...
10            # ...
```


has same shape as input

scalar



Computational graph

Reductive ops — sum



```
1 def sum(input,axis = None, keepdims = False):
2     output = ag.Tensor(np.sum(input.value, axis = axis, keepdims =
3         keepdims), inputs = [input], op='sum')
4     def _backward():
5         if axis == None or keepdims:
6             input.grad += output.grad
7         else:
8             input.grad += np.expand_dims(output.grad, axis = axis)
9         return None
10 # ...
11 # ...
```

← broad casting

Exercises

- Reshape (Exercise 2.1)
 - Fix the backward function for reshape

Solution

```
1  def reshape(input, newshape):
2      output = ag.Tensor(np.reshape(input.value, newshape), inputs=[
input], op="reshape")
3      def _backward():
4          input.grad += np.reshape(output.grad, input.shape)
5          return None
6      output._backward = _backward
7      return output
```

Handwritten annotations:

- A bracket under `input.value` with the label `oldshape`.
- A bracket under `input.shape` with the label `oldshape`.

Exercises

- Reshape (Exercise 2.1)
 - Fix the backward function for reshape
- Indexing (Exercise 2.2)
 - Fix the loss derivative autograd calculation

Exercises

- Reshape (Exercise 2.1)
 - Fix the backward function for reshape
- Indexing (Exercise 2.2)
 - Fix the loss derivative autograd calculation
- Implementing BCE (Exercise 2.3)
 - Implement the binary cross entropy.

Convolution

Class 2: Bird




Class 1: Automobile




$X.shape = (N, C, H, W) = (2, 3, 32, 32)$ (2 images from CIFAR10)

Handwritten annotations:
- A bracket under the first '2' in the tuple points to the text "2 images".
- A bracket under the '3' points to the text "RGB channels".
- A bracket under the first '32' points to the text "height".
- A bracket under the second '32' points to the text "width".

Convolution




$$\mathbf{X}.\text{shape} = (N, C, H, W) = (2, 3, 32, 32)$$

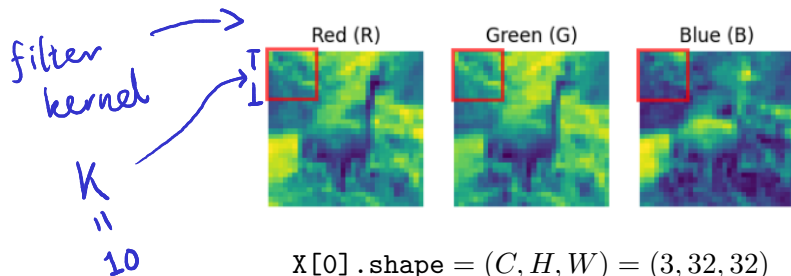


$$\mathbf{X}[0] . \text{shape} = (C, H, W) = (3, 32, 32)$$


Convolution




$X.shape = (N, C, H, W) = (2, 3, 32, 32)$



Convolution




$$\mathbf{X}.\text{shape} = (N, C, H, W) = (2, 3, 32, 32)$$




$$\mathbf{X}[0].\text{shape} = (C, H, W) = (3, 32, 32)$$

Convolution




$$X.shape = (N, C, H, W) = (2, 3, 32, 32)$$




$$X[0].shape = (C, H, W) = (3, 32, 32)$$

Convolution



$$X.shape = (N, C, H, W) = (2, 3, 32, 32)$$



$$X[0].shape = (C, H, W) = (3, 32, 32)$$

Convolution

Restrict ourselves to kernel size as the only hyperparam

Class 2: Bird




Class 1: Automobile




$$X.shape = (N, C, H, W) = (2, 3, 32, 32)$$


Red (R)



Green (G)



Blue (B)



move 1 step at time
stride = 1
pad = 0

$$\begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

*

$$\begin{bmatrix} ? & ? \\ ? & ? \end{bmatrix}$$

$$X[0].shape = (C, H, W) = (3, 32, 32)$$

multiply entrywise

sum over pixels → sum over channels

LeNet5




Image from LeCun et al. 1998


Convolution

allows you to leverage
still faster
leverage (C, W) BLAS \rightarrow linear algebra

"im2col"

\gg for loops

no additional
space requirement



flattened
patch

$X[0].\text{shape} = (C, H, W) = (3, 32, 32)$

flatten

patches

Every images gets
pulled into a
stack of
patches



$X_patches[0].\text{shape} = (P, \underline{C}, \underline{K}, \underline{K}) = (3, 32, 32)$

K = size of window, P = number of patches


Best conv
Intel MKL
Cuda NVIDIA

Convolution

Stanford cs231n



$X[0].\text{shape} = (C, H, W) = (3, 32, 32)$



$X_{\text{patches}}[0].\text{shape} = (P, C, K, K) = (3, 32, 32)$

K = size of window, P = number of patches

A simple example

im2col

high level
requires

matmul

reshape

ch H W

```
1 C, H, W = 2, 4, 5
2 K = 2 # Kernel size
3 N = 2 # Batch size
4 X = np.arange(N * C * H * W).reshape(N, C, H, W)
5 X[0]
6 array([[ 0,  1,  2,  3,  4],
7        [ 5,  6,  7,  8,  9],
8        [10, 11, 12, 13, 14],
9        [15, 16, 17, 18, 19]],
10
11        [[20, 21, 22, 23, 24],
12         [25, 26, 27, 28, 29],
13         [30, 31, 32, 33, 34],
14         [35, 36, 37, 38, 39]])
```

← channel 0

← channel 1

Goal: compute X_patches. Okay, but what is X_patches?

A simple example

$$(H-1)(W-1)$$

```
1 X[0]
2 array([[ 0,  1,  2,  3,  4],
3        [ 5,  6,  7,  8,  9],
4        [10, 11, 12, 13, 14],
5        [15, 16, 17, 18, 19]],
6
7        [[20, 21, 22, 23, 24],
8         [25, 26, 27, 28, 29],
9         [30, 31, 32, 33, 34],
10        [35, 36, 37, 38, 39]])
```

$$X_{\text{patches}}. \text{shape} = [N, P, C, K, K]$$

↑
number of patches

Goal: compute `X_patches`. Okay, but what is `X_patches`?

```
1 X_patches[0,0]
2 array([[ 0.,  1.],
3        [ 5.,  6.]],
4
5        [[20., 21.],
6         [25., 26.]])
```

A simple example

```
1 X[0]
2 array([[ 0,  1,  2,  3,  4],
3        [ 5,  6,  7,  8,  9],
4        [10, 11, 12, 13, 14],
5        [15, 16, 17, 18, 19]],
6
7        [[20, 21, 22, 23, 24],
8         [25, 26, 27, 28, 29],
9         [30, 31, 32, 33, 34],
10        [35, 36, 37, 38, 39]])
```

Goal: compute `X_patches`. Okay, but what is `X_patches`?

```
1 X_patches[0,1]
2 array([[ 1.,  2.],
3        [ 6.,  7.]],
4
5        [[21., 22.],
6         [26., 27.]])
```

patch index

A simple example

```
1 X[0]
2 array([[ 0,  1,  2,  3,  4],
3        [ 5,  6,  7,  8,  9],
4        [10, 11, 12, 13, 14],
5        [15, 16, 17, 18, 19]],
6
7        [[20, 21, 22, 23, 24],
8         [25, 26, 27, 28, 29],
9         [30, 31, 32, 33, 34],
10        [35, 36, 37, 38, 39]])
```

Goal: compute `X_patches`. Okay, but what is `X_patches`?

```
1 X_patches[0,3]
2 array([[ 3.,  4.],
3        [ 8.,  9.]],
4
5        [[23., 24.],
6         [28., 29.]])
```

A simple example

```
1 X[0]
2 array([[ 0,  1,  2,  3,  4],
3        [ 5,  6,  7,  8,  9],
4        [10, 11, 12, 13, 14],
5        [15, 16, 17, 18, 19]],
6
7        [[20, 21, 22, 23, 24],
8        [25, 26, 27, 28, 29],
9        [30, 31, 32, 33, 34],
10       [35, 36, 37, 38, 39]])
```

Goal: compute `X_patches`. Okay, but what is `X_patches`?

```
1 X_patches[0,4]
2 array([[ 5.,  6.],
3        [10., 11.]],
4
5        [[25., 26.],
6        [30., 31.]])
```

A simple example

```
1 X[1]
2 array([[40, 41, 42, 43, 44],
3        [45, 46, 47, 48, 49],
4        [50, 51, 52, 53, 54],
5        [55, 56, 57, 58, 59]],
6
7        [[60, 61, 62, 63, 64],
8        [65, 66, 67, 68, 69],
9        [70, 71, 72, 73, 74],
10       [75, 76, 77, 78, 79]])
```

Goal: compute `X_patches`. Okay, but what is `X_patches`?

```
1 X_patches[1,0]
2 array([[40., 41.],
3        [45., 46.]],
4
5        [[60., 61.],
6        [65., 66.]])
```

X_patches

$$\text{X_patches}[0].\text{shape} = (P, C, K, K) = (3, 32, 32)$$

K = size of window, P = number of patches

```
1 CHW = C * H * W
2 out_H = H - K + 1 # Output height
3 out_W = W - K + 1 # Output width
4 P = out_H * out_W # Total number of patches per image
```

Strategy:

- First flatten X to be a matrix, call it X_{flat}

X_patches

$$\text{X_patches}[0].\text{shape} = (P, C, K, K) = (3, 32, 32)$$

K = size of window, P = number of patches

```
1 CHW = C * H * W
2 out_H = H - K + 1 # Output height
3 out_W = W - K + 1 # Output width
4 P = out_H * out_W # Total number of patches per image
```

Strategy:

- First flatten X to be a matrix, call it X_{flat}
- Pick out the patches using this “im2col_mat” matrix

X_patches

$$\text{X_patches}[0].\text{shape} = (P, C, K, K) = (3, 32, 32)$$

K = size of window, P = number of patches

```
1 CHW = C * H * W
2 out_H = H - K + 1 # Output height
3 out_W = W - K + 1 # Output width
4 P = out_H * out_W # Total number of patches per image
```

Strategy:

- First flatten X to be a matrix, call it X_{flat}
- Pick out the patches using this “`im2col_mat`” matrix
- (there are two flavors: dense and sparse. More on that later)

X_patches

$$\text{X_patches}[0].\text{shape} = (P, C, K, K) = (3, 32, 32)$$

K = size of window, P = number of patches


```
1 CHW = C * H * W
2 out_H = H - K + 1 # Output height
3 out_W = W - K + 1 # Output width
4 P = out_H * out_W # Total number of patches per image
```

Strategy:

- First flatten X to be a matrix, call it X_{flat}
- Pick out the patches using this “`im2col_mat`” matrix
- (there are two flavors: dense and sparse. More on that later)
- Get $X_{\text{patches_flat}}$. Reshape them to get X_{patches} .

A simple example

```
1 C, H, W = 2, 4, 5
2 K,N = 2,2 # Kernel size, Batch size
3 X = np.arange(N * C * H * W).reshape(N, C, H, W)
4 X[0]
5 array([[[ 0,  1,  2,  3,  4],
6         [ 5,  6,  7,  8,  9],
7         [10, 11, 12, 13, 14],
8         [15, 16, 17, 18, 19]],
9
10        [[20, 21, 22, 23, 24],
11        # ...
12 X_flat = X.reshape(N, -1) # Shape (N, C*H*W)
```



A simple example

```
1 C, H, W = 2, 4, 5
2 K,N = 2,2 # Kernel size, Batch size
3 X = np.arange(N * C * H * W).reshape(N, C, H, W)
4 X[0]
5 array([[0, 1, 2, 3, 4],
6        [5, 6, 7, 8, 9],
7        [10, 11, 12, 13, 14],
8        [15, 16, 17, 18, 19]],
9
10       [[20, 21], 22, 23, 24],
11 # ...
12 X_flat = X.reshape(N, -1) # Shape (N, C*H*W)
13 X_flat[0]
14 array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15,
15        16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32,
16        33, 34, 35, 36, 37, 38, 39])
```

A simple example

```
1 X_flat = X.reshape(N, -1) # Shape (N, C*H*W)
2 X_flat[0]
3 array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15,
        16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32,
        33, 34, 35, 36, 37, 38, 39])
4
5 im2col_mat_dense = im2col_matrix_dense(X, K) # assume this given
6 X_out_dense = np.matmul(X_flat, im2col_mat_dense)
7 X_out_dense
8 array([[0., 1., 5., 6., 20., 21., 25., 26., 1., 2., 6., 7., 21.,
        22., 26., 27., 2., 3., 7., 8., 22., 23., 27., 28., 3., 4.,
        8., 9., 23., 24., 28., 29., 5., 6., 10., 11., 25., 26., 30.,
        31., 6., 7., 11., 12., 26., 27., 31., 32., 7., 8., 12., 13.,
        ...])
```

A simple example

```
1 X_out_dense = np.matmul(X_flat, im2col_mat_dense)
2 X_out_dense
3 array([[ 0.,  1.,  5.,  6., 20., 21., 25., 26.,  1.,  2.,  6.,  7., 21.,
4         22., 26., 27.,  2.,  3.,  7.,  8., 22., 23., 27., 28.,  3.,  4.,
5         8.,  9., 23., 24., 28., 29.,  5.,  6., 10., 11., 25., 26., 30.,
6         31.,  6.,  7., 11., 12., 26., 27., 31., 32.,  7.,  8., 12., 13.,
7         ...
8 X_patches = X_patches_flat.reshape(N,P,C,K,K)
```

A simple example

```
1 X_out_dense = np.matmul(X_flat, im2col_mat_dense)
2 X_out_dense
3 array([[ 0.,  1.,  5.,  6., 20., 21., 25., 26.,  1.,  2.,  6.,  7., 21.,
4         22., 26., 27.,  2.,  3.,  7.,  8., 22., 23., 27., 28.,  3.,  4.,
5         8.,  9., 23., 24., 28., 29.,  5.,  6., 10., 11., 25., 26., 30.,
6         31.,  6.,  7., 11., 12., 26., 27., 31., 32.,  7.,  8., 12., 13.,
7         ...
8 X_patches = X_patches_flat.reshape(N,P,C,K,K)
9 X_patches[0,0]
10 array([[ 0.,  1.],
11         [ 5.,  6.]],
12
13         [[20., 21.],
14         [25., 26.]])
```

Exercise 3


-1x zero everywhere else

```
1  # input
2  X = np.arange(N * C * H * W).reshape(N, C, H, W)
3  X[0]
4  array([[[ 0,  1,  2,  3,  4],
5          [ 5,  6,  7,  8,  9],
6          [10, 11, 12, 13, 14],
7          [15, 16, 17, 18, 19]],
8         # ...
9
10         # desired output
11         # (2 times lower left - 1 times top right, ignore second channel)
12  X_convolved.shape = (2, 3, 4)
13  X_convolved[0]
14  array([[ 9., 10., 11., 12.],
15          [14., 15., 16., 17.],
16          [19., 20., 21., 22.]])
```

Handwritten annotations:


- A red box highlights the 2x2 sub-region of the input array containing values 0, 1, 5, and 6.
- A blue box highlights the 2x2 sub-region of the input array containing values 6, 7, 11, and 12.
- A red '2x' is written next to the first two rows of the input array.
- A blue box highlights the value 15 in the output array.
- A red circle highlights the value 9 in the output array.

Smoothness



```
1 Conv_kernel = np.array([[[ 0, -1.], [ 1., 0.]], [[0., 0.], [0., 0.]],  
    [[0., 0.], [0., 0.]])
```

Smoothness



```
1 Conv_kernel = np.array([[[ 1, 0.], [ 0., -1.]], [[0., 0.], [0., 0.]],  
    [[0., 0.], [0., 0.]])
```

References I

- [LeC+98] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. “Gradient-based learning applied to document recognition”. In: *Proceedings of the IEEE* 86.11 (1998), pp. 2278–2324.