

More convolutional nets


Lecture 09 — CS 577 Deep Learning

Instructor: Yutong Wang

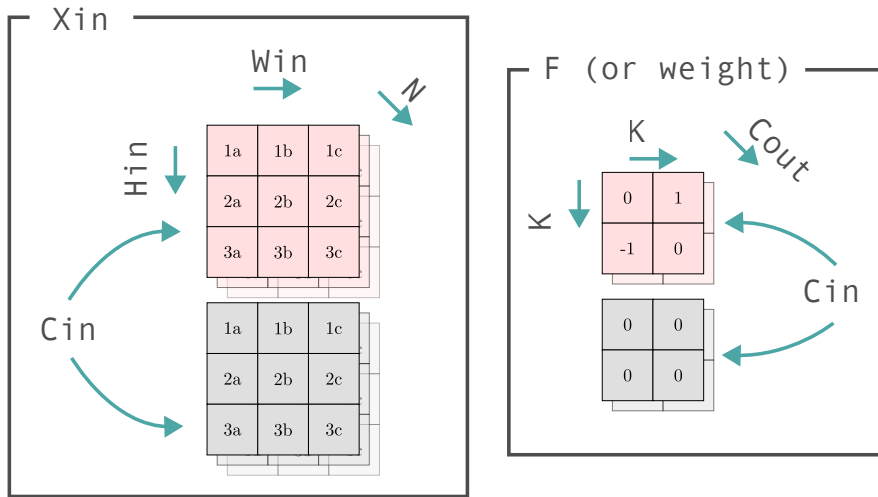
Computer Science
Illinois Institute of Technology

October 16, 2024

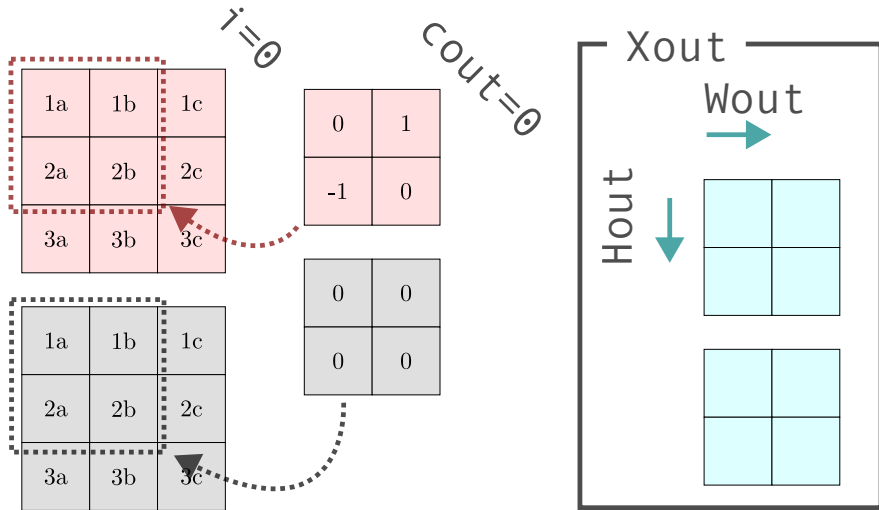
Input image tensor with shape $(N, C_{\text{in}}, H_{\text{in}}, W_{\text{in}})$




A single convolution layer $(C_{\text{out}}, C_{\text{in}}, K, K)$



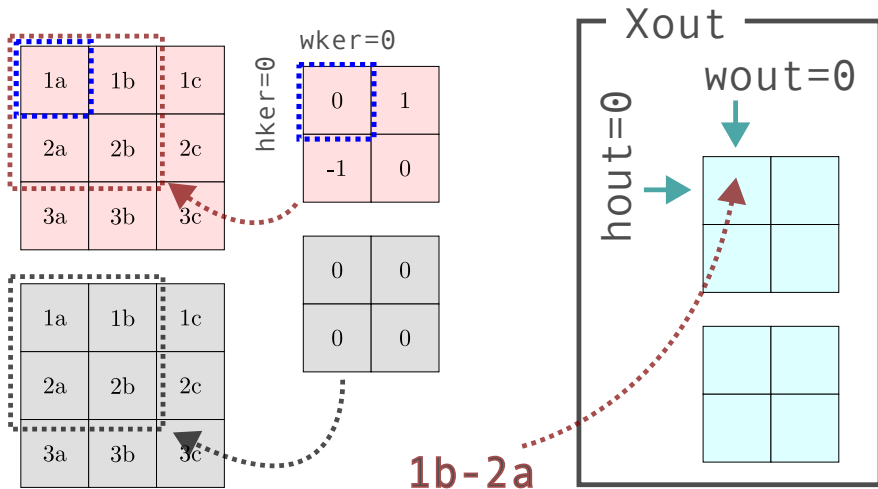
Output: Xout with shape $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$



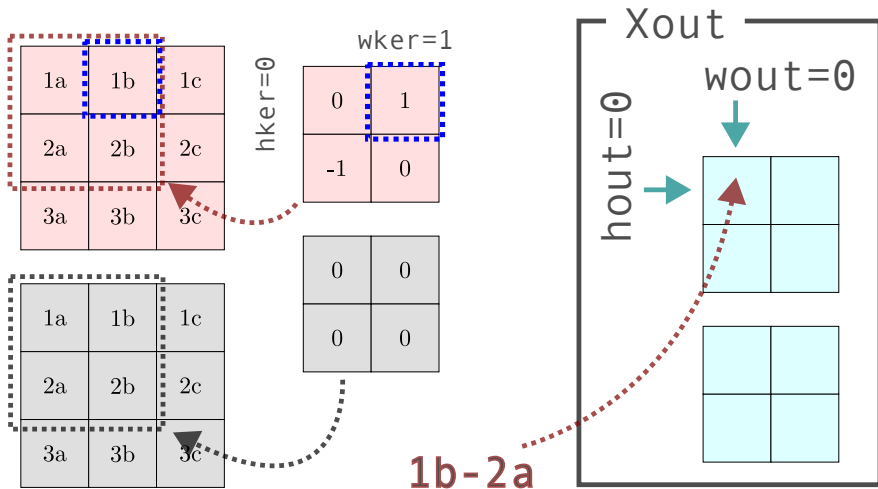
patch_idx = 0




patch_idx = 0



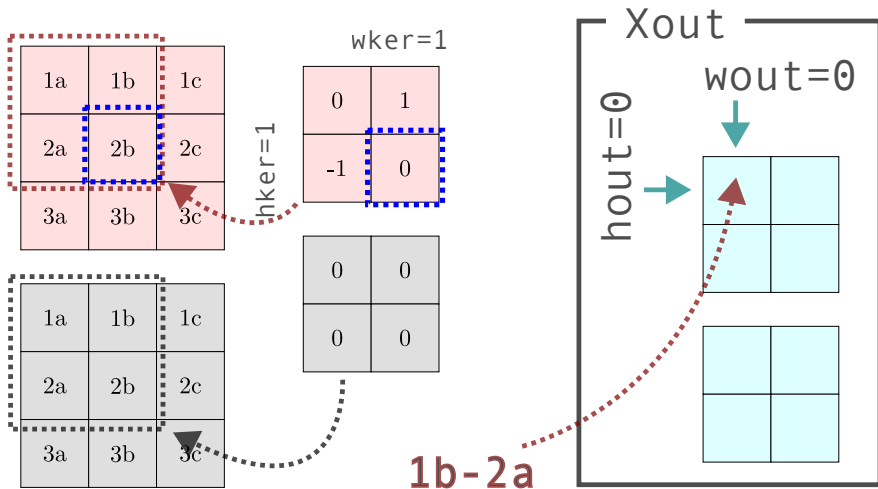
patch_idx = 0




patch_idx = 0



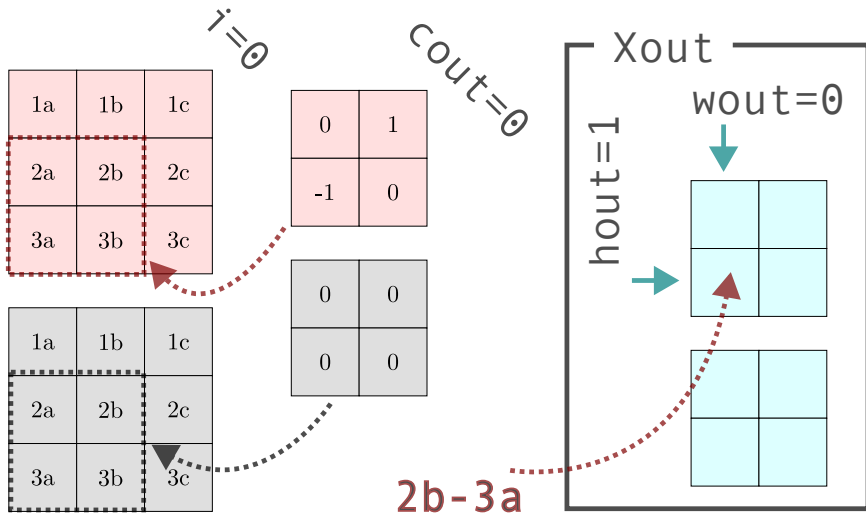
patch_idx = 0



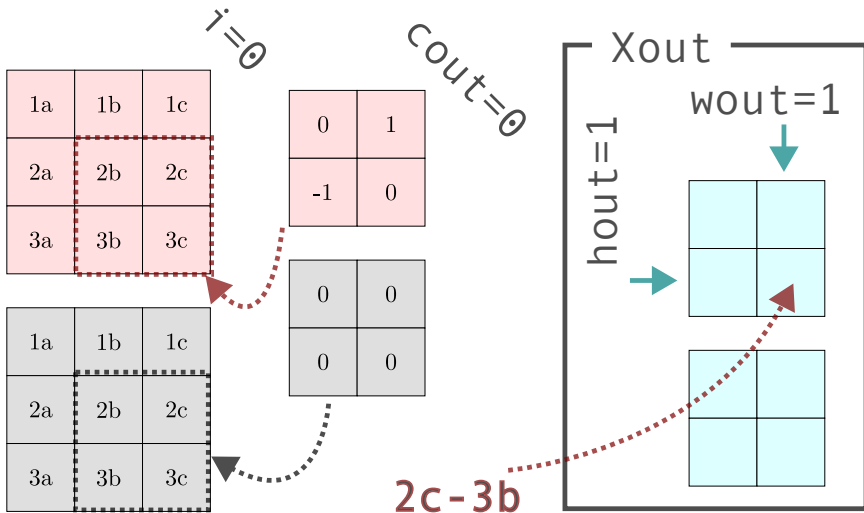
patch_idx = 1



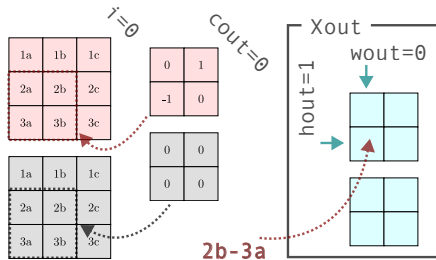
patch_idx = 2



patch_idx = 3



Sliding window (simple but slow)



```
1 # Suppose i, c_out, h_out, w_out are already defined, and
2 # X_out is initialized to all zeros
3 for cin in range(Cin): # Input channels
4     for hker in range(K): # Kernel height
5         for wker in range(K): # Kernel width
6             Xout[i, cout, hout, wout] += (
7                 F[cout, cin, hker, wker] *
8                 Xin[i, cin, hout + hker, wout + wker])
```

There's gotta be a better way! "im2col" (Next)

“im2col”

Main idea (in theory)

- Keep the window fixed in one spot
- Move the image tensor

Main idea (in theory)

- Create patches of the data
- Each patch has the same shape as the convolution filter
- Do matrix multiplication between the (flattened) patches and the (flattened) convolutional filter weights

Flatten the weights

1a	1b	1c
2a	2b	2c
3a	3b	3c

1a	1b	1c
2a	2b	2c
3a	3b	3c

0	1
-1	0

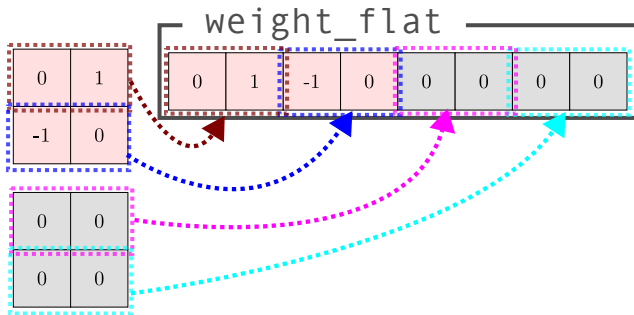
0	0
0	0

weight_flat

Flatten the weights

1a	1b	1c
2a	2b	2c
3a	3b	3c

1a	1b	1c
2a	2b	2c
3a	3b	3c



1a	1b	1c
2a	2b	2c
3a	3b	3c


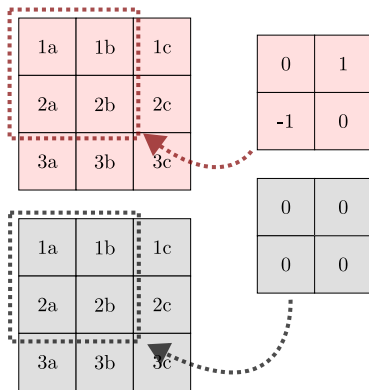
0	0
0	0

weight_flat


0	1	-1	0	0	0	0	0
---	---	----	---	---	---	---	---

Xin_patches_flat


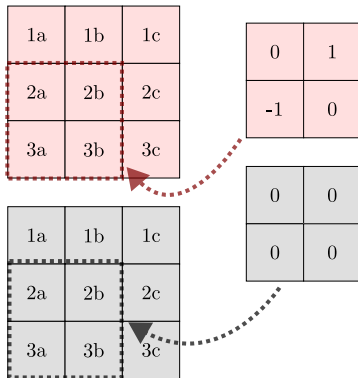
Flatten the weights



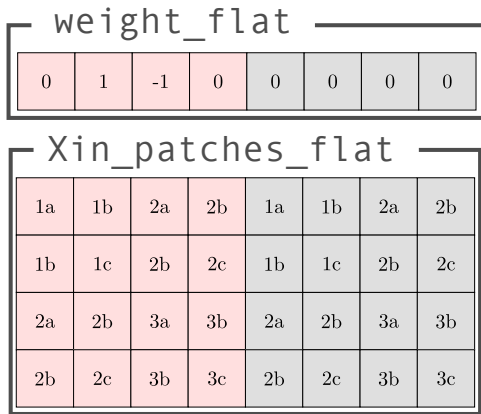

Flatten the weights



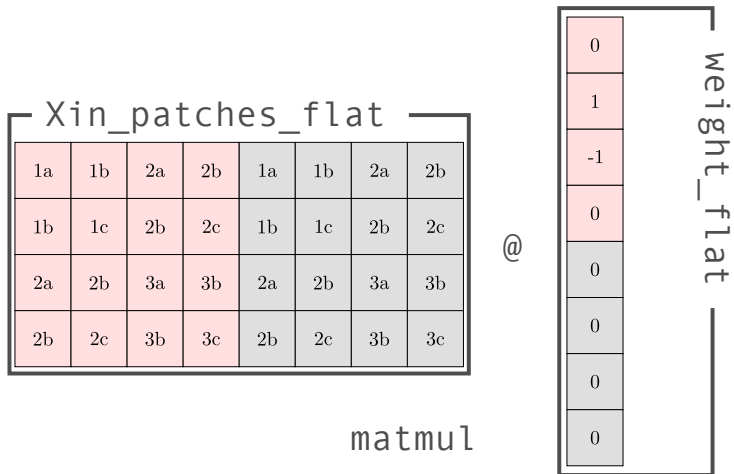
Flatten the weights




Flatten the weights




Sliding window via matmul



Sliding window via matmul



Flattening Xin into Xin_flat



Flattening Xin into Xin_flat

Recall patch_idx=0:

1a	1b	1c
2a	2b	2c
3a	3b	3c

1a	1b	1c
2a	2b	2c
3a	3b	3c

Our goal:

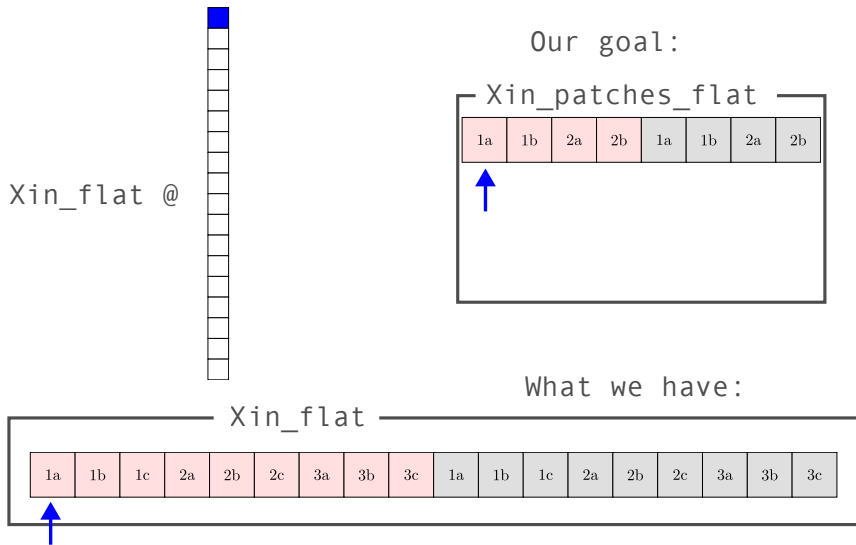
Xin_patches_flat							
1a	1b	2a	2b	1a	1b	2a	2b

What we have:

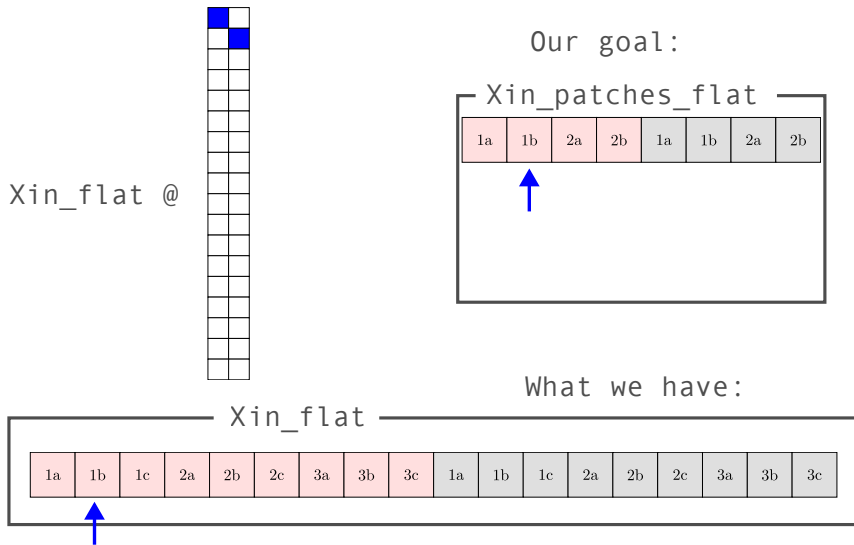
Xin_flat

1a	1b	1c	2a	2b	2c	3a	3b	3c	1a	1b	1c	2a	2b	2c	3a	3b	3c
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

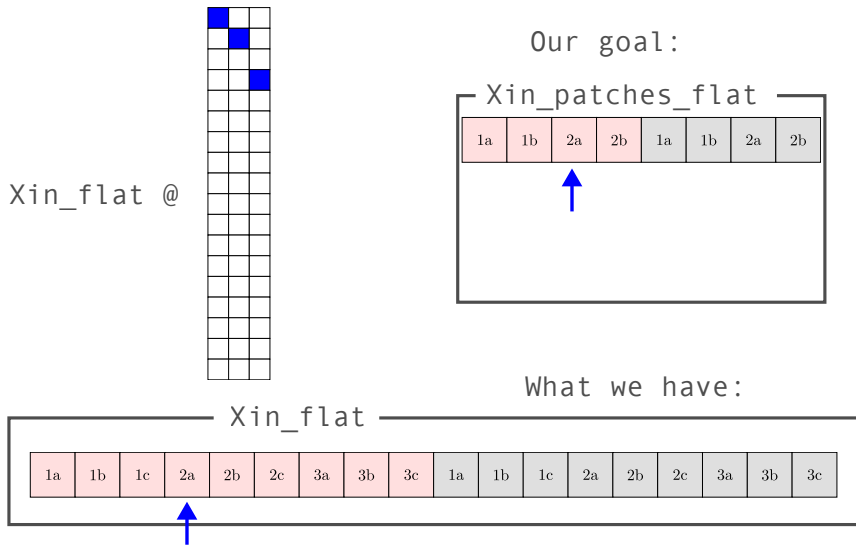
The “im2col” matrix




The “im2col” matrix




The “im2col” matrix




The “im2col” matrix



Our goal:

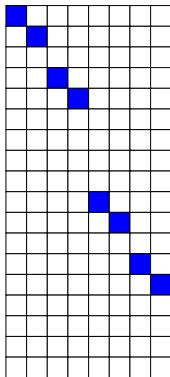


What we have:

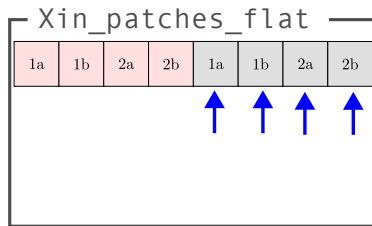


The “im2col” matrix

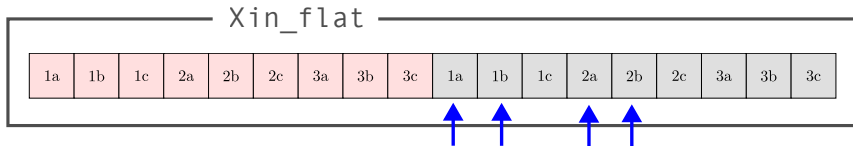
Xin_flat @



Our goal:



What we have:



The “im2col” matrix

Recall `patch_idx=1`:

1a	1b	1c
2a	2b	2c
3a	3b	3c

1a	1b	1c
2a	2b	2c
3a	3b	3c

Our goal:

`Xin_patches_flat`

1b	1c	2b	2c	1b	1c	2b	2c
----	----	----	----	----	----	----	----

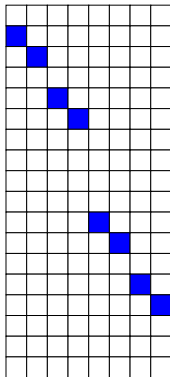
What we have:

`Xin_flat`

1a	1b	1c	2a	2b	2c	3a	3b	3c	1a	1b	1c	2a	2b	2c	3a	3b	3c
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

The “im2col” matrix

Xin_flat @



Our goal:

Xin_patches_flat

1b	1c	2b	2c	1b	1c	2b	2c
----	----	----	----	----	----	----	----

What we have:

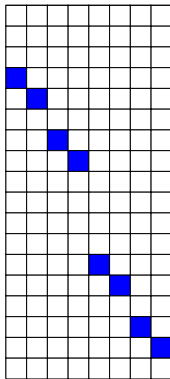
Xin_flat

1a	1b	1c	2a	2b	2c	3a	3b	3c	1a	1b	1c	2a	2b	2c	3a	3b	3c
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----



The “im2col” matrix

Xin_flat @




Our goal:

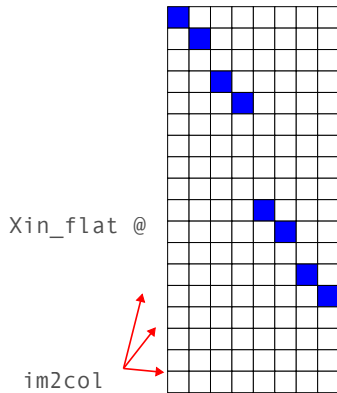
Xin_patches_flat

2a	2b	3a	3b	2a	2b	3a	3b
----	----	----	----	----	----	----	----

What we have:



The “im2col” matrix




=

Xin_im2col

1a	1b	2a	2b	1a	1b	2a	2b
----	----	----	----	----	----	----	----

The “im2col” matrix




=

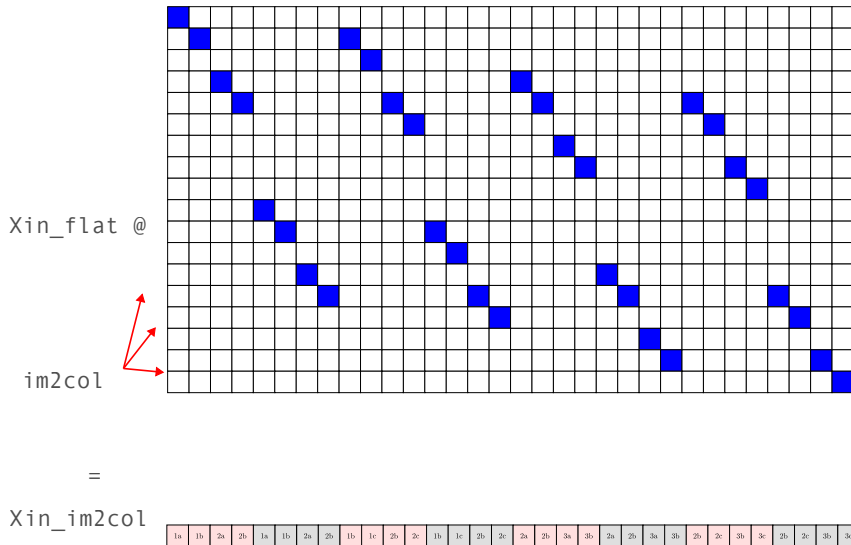
Xin_im2col

1a	1b	2a	2b	1a	1b	2a	2b	1b	1c	2b	2c	1b	1c	2b	2c
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

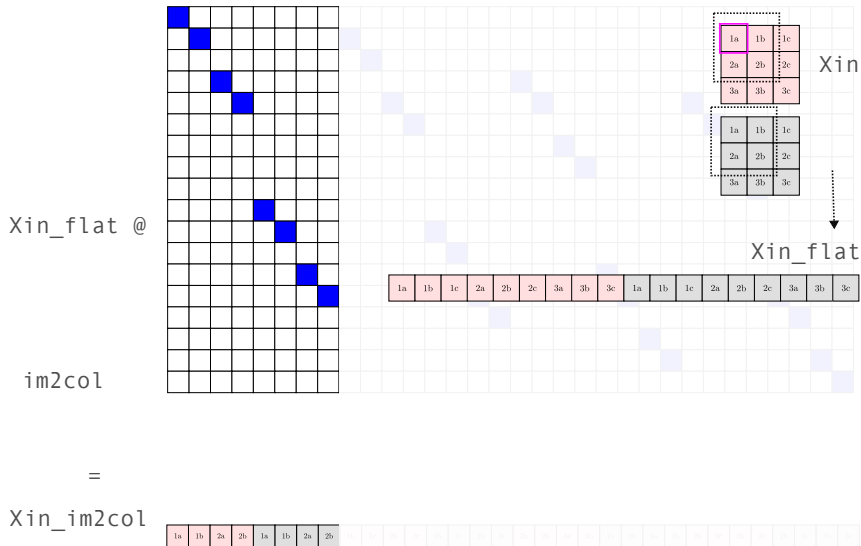
The “im2col” matrix




The “im2col” matrix



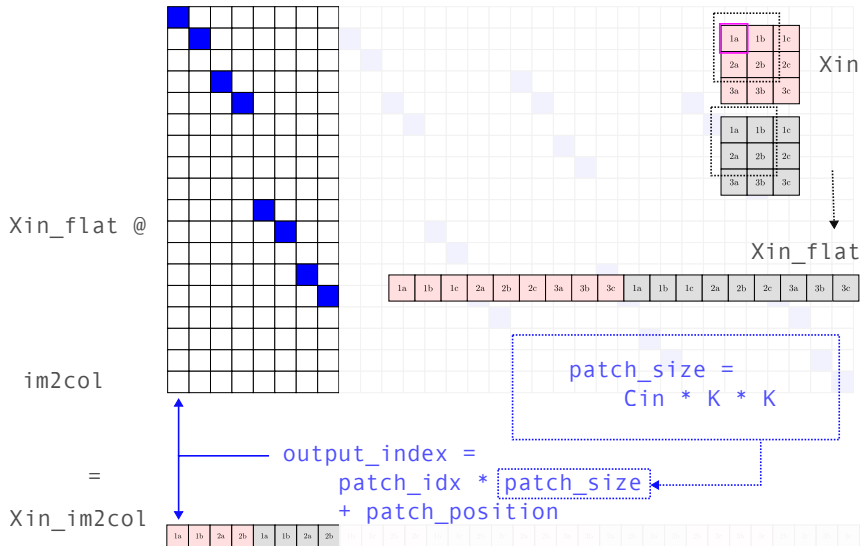
The “im2col” matrix



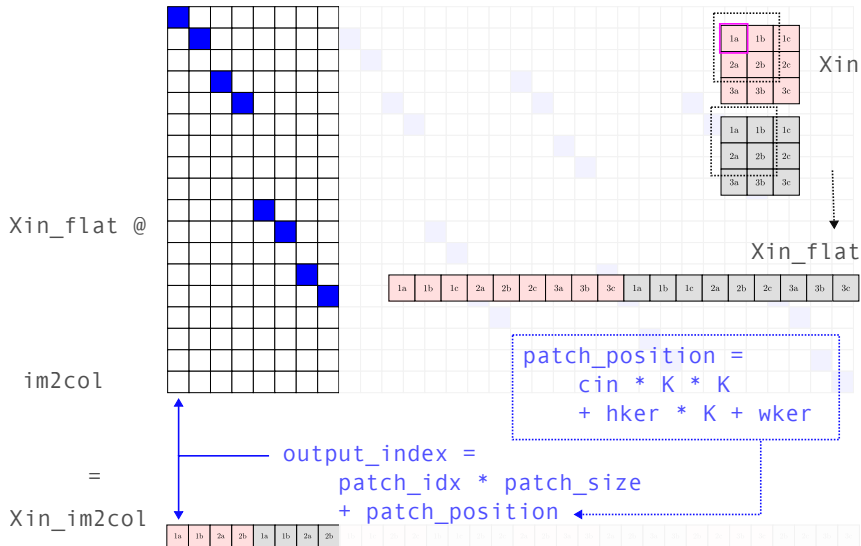
The “im2col” matrix




The “im2col” matrix



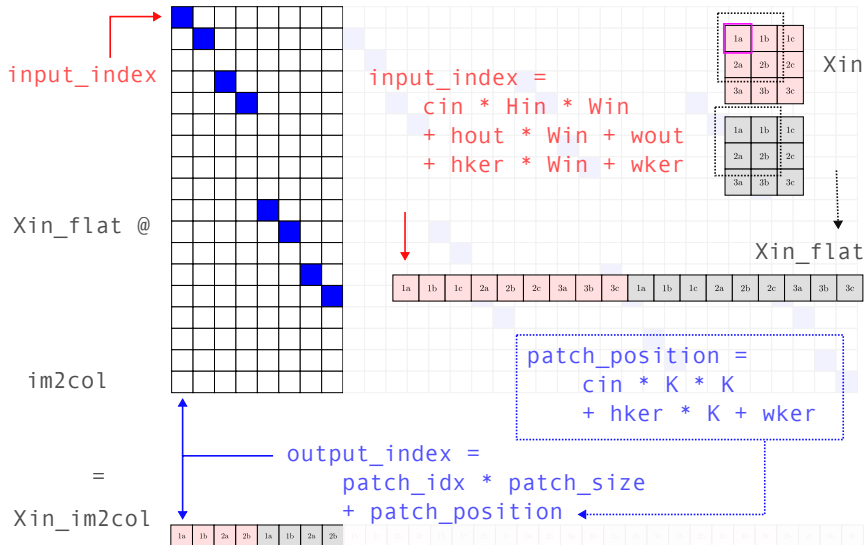
The “im2col” matrix




The “im2col” matrix




The “im2col” matrix




The “im2col” matrix



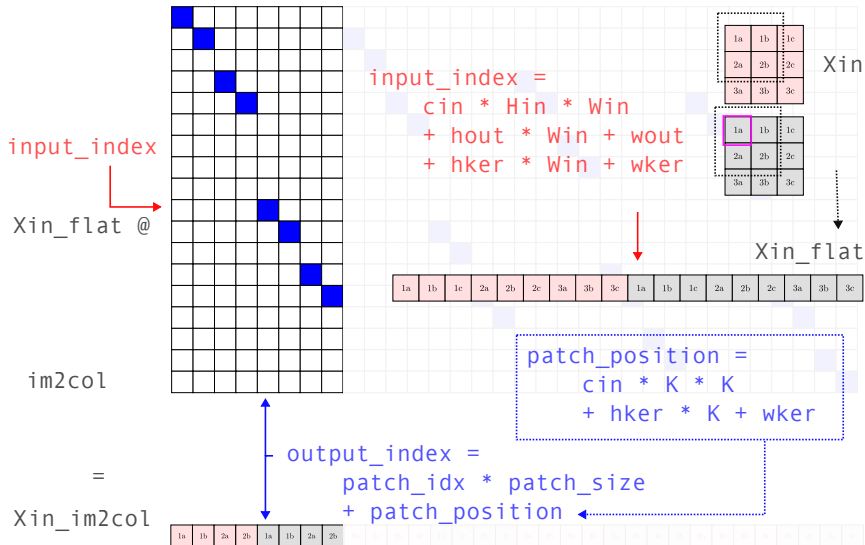
The “im2col” matrix




The “im2col” matrix




The “im2col” matrix



The “im2col” matrix



The “im2col” matrix



Putting it all together

```
1 def im2col_matrix_dense(Xin, K, S=1):
2     N, Cin, Hin, Win = Xin.shape
3     Hout, Wout = Hin - K + 1, Win - K + 1
4     P = Hout * Wout # Total number of patches per image
5     patch_size = Cin * K * K # Size of each flattened patch
6     im2col_mat_dense = np.zeros((Cin * Hin * Win, P * patch_size))
7
8     # [main loop on next slide...]
9
10    return im2col_mat_dense
```

Putting it all together

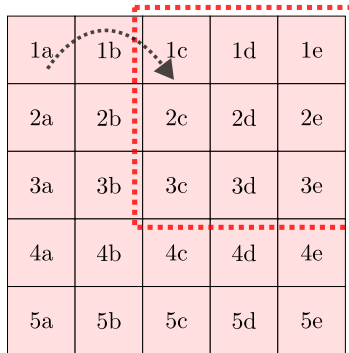
```
1  # [continued from previous slide...]
2  patch_idx = 0
3  for hout in range(Hout):
4      for wout in range(Wout):
5          for cin in range(Cin):
6              for hker in range(K):
7                  for wker in range(K):
8                      input_index = cin * Hin * Win + hout * Win +
wout + hker * Win + wker
9                      patch_position = cin * K * K + hker * K + wker
10                     output_index = patch_idx * patch_size +
patch_position
11                     im2col_mat_dense[input_index, output_index] = 1
12                     patch_idx += 1
```

Stride = 2

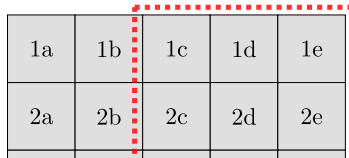
1a	1b	1c	1d	1e
2a	2b	2c	2d	2e
3a	3b	3c	3d	3e
4a	4b	4c	4d	4e
5a	5b	5c	5d	5e

1a	1b	1c	1d	1e
2a	2b	2c	2d	2e

Stride = 2

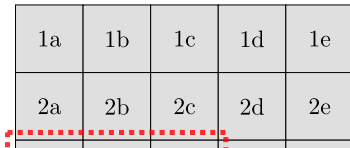
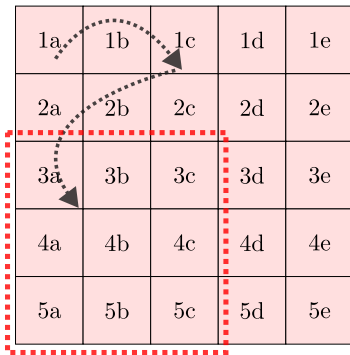


1a	1b	1c	1d	1e
2a	2b	2c	2d	2e
3a	3b	3c	3d	3e
4a	4b	4c	4d	4e
5a	5b	5c	5d	5e



1a	1b	1c	1d	1e
2a	2b	2c	2d	2e

Stride = 2



Implement `im2col_matrix_dense` with stride

See `lec09-in-class-ex1-im2col.ipynb`

End of part 1.

Goal for today: implement LeNet5





Image from LeCun et al. 1998

$$\text{im2col.shape} = C_{\text{in}} \cdot H_{\text{in}} \cdot W_{\text{in-by-}P} \cdot C_{\text{out}} \cdot K^2$$



$$\text{im2col.shape} = C_{\text{in}} \cdot H_{\text{in}} \cdot W_{\text{in}}\text{-by-}P \cdot C_{\text{in}} \cdot K^2$$

For 1st convolution layer in LeNet5.

- $H_{\text{in}} = W_{\text{in}} = 32$
- $C_{\text{in}} = 1$
- $K = 5$
- $S = 1$
- $H_{\text{out}} = H_{\text{in}} - K + 1 = 28$
- $W_{\text{out}} = W_{\text{in}} - K + 1 = 28$
- $P = 28 * 28$
- Final shape of `im2col` = 1024 -by- 19600
- 20,070,400 float64's to store \approx 160 MB

Sparse matrices


Pros

- Store just the non-zero values of the matrix (memory efficient)
- Faster than dense matmul when matrices are mostly zeros (speedier)

Cons

- Irregular memory access (less cache-optimization friendly)
- Nightmare for GPUs (as of now, to the best of my knowledge)

from scipy.sparse import csr_matrix



```
1 data = [1,1,1,1,1,...] # len(data) = P * patch_size
2 row_indices = [0,1,3,4,9,...]
3 col_indices = [0,1,2,3,4,...]
4 im2col_mat_sparse = csr_matrix((data,
5                                 (row_indices, col_indices)),
6                                 shape=(n_rows, n_cols))
```

Example

```
1 data = [1,2,3,1]
2 row_indices = [1,2,0,1]
3 col_indices = [0,1,2,1]
4 sparse_mat_example = csr_matrix((data, (row_indices, col_indices)),
    shape=(3, 3))
5 sparse_mat_example.toarray()
```

Output:

```
1 array([[0, 0, 3],
2        [1, 1, 0],
3        [0, 2, 0]])
```

Dense-sparse multiplication

```
1 X = np.arange(9).reshape(3,3)
2 ## this won't work
3 # np.matmul(X, sparse_mat_example)
4 ## this works bc method resolution order
5 X @ sparse_mat_example
```

```
1 array([[0, 1, 2],      array([[0, 0, 3],      array([[ 1,  5,  0],
2         [3, 4, 5],      [1, 1, 0],      =         [ 4, 14,  9],
3         [6, 7, 8]])     [0, 2, 0]])     [ 7, 23, 18]])
```

Implement `im2col_matrix_sparse`

- `lec09-in-class-ex2-sparse.ipynb`

Quick update

Can add/mul constants both on left and right now.

```
1     def __radd__(self, other):
2         return self + other
3
4     def __rmul__(self, other):
5         return self * other
6
7     def __rsub__(self, other):
8         return (-self) + other
9
10    def __rtruediv__(self, other):
11        return ag.Tensor(other) / self
```


Quick update

```
1  def moveaxis(input, source, destination):
2      output = ag.Tensor(np.moveaxis(input.value, source, destination)
3      , inputs=[input], op="moveaxis")
4
5      def _backward():
6          input.grad += np.moveaxis(output.grad, source, destination)
7          return None
8      output._backward = _backward
9      return output
```


Enabling autograd for dense-sparse matmul

- Problem: our current autograd-enabled matmul only supports when both inputs are numpy arrays
- Observation: we don't need to track gradients on the `im2col_mat`
- Solution:

```
1  # [sp]arse [c]onstant (non-AG-enabled) [mat]rix [mul]tiplication
2  def spcmatmul(input, sparse_mat: csr_matrix):
3      output = ag.Tensor(input.value @ sparse_mat,
4                          inputs = [input],
5                          op="spcmatmul")
6
7      def _backward():
8          input.grad += output.grad @ sparse_mat.T
9          return None
10
11     output._backward = _backward
12     return output
```


Exercise 3

- Create the Conv2d layer
- Create the AvgPool2d layer
- Create the LeNet5 model




- Create the `nn.CrossEntropyLoss`

Sanity check for conv2d



Sanity check for pool



Sanity check for training run

You should achieve 100 percentage training accuracy

References I

- [LeC+98] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. “Gradient-based learning applied to document recognition”. In: *Proceedings of the IEEE* 86.11 (1998), pp. 2278–2324.