

More convolutional nets


Lecture 09 — CS 577 Deep Learning

Instructor: Yutong Wang

Computer Science
Illinois Institute of Technology

October 16, 2024


Input image tensor with shape $(N, C_{in}, H_{in}, W_{in})$




Samples
n channels-in
1 grayscale
3 RGB
height width
32 x 32
padded mnist


A single convolution layer (C_{out}, C_{in}, K, K)

abstract \downarrow *physical*




Output: Xout with shape $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$







patch_idx = 0




patch_idx = 0



patch_idx = 0




patch_idx = 0




`patch_idx = 1`


"1c" entire symbol substitute for actual number




patch_idx = 2



patch_idx = 3



Sliding window (simple but slow)



Full conv
has
7 nested
for-loops

```
1 # Suppose i, c_out, h_out, w_out are already defined, and
```

```
2 # X_out is initialized to all zeros
```

```
3 for cin in range(Cin): # Input channels
```

```
4     for hker in range(K): # Kernel height
```

```
5         for wker in range(K): # Kernel width
```

```
6             Xout[i, c_out, hout, wout] += (
```


```
7                 F[c_out, cin, hker, wker] *
```

```
8                 Xin[i, cin, hout + hker, wout + wker])
```

patch-idx

There's gotta be a better way! "im2col" (Next)

Sliding window (simple but slow)



```
1 # Suppose i, c_out, h_out, w_out are already defined, and
2 # X_out is initialized to all zeros
3 for cin in range(Cin): # Input channels
4     for hker in range(K): # Kernel height
5         for wker in range(K): # Kernel width
6             Xout[i, c_out, hout, wout] += (
7                 F[cout, cin, hker, wker] *
8                 Xin[i, cin, hout + hker, wout + wker])
```

im2col

There's gotta be a better way! "im2col" (Next)

"im2col"



Conv-Filter

Main idea (in theory)

- Keep the window fixed in one spot
- Move the image tensor

Main idea (in theory) practice

- Create patches of the data
- Each patch has the same shape as the convolution filter
- Do matrix multiplication between the (flattened) patches and the (flattened) convolutional filter weights



"lower the
convolution
to matmul"

“im2col”

Main idea (in theory)

- Keep the window fixed in one spot
- Move the image tensor

Main idea (in theory)

- Create patches of the data
- Each patch has the same shape as the convolution filter
- Do matrix multiplication between the (flattened) patches and the (flattened) convolutional filter weights

im2col physical interpretation
many way to compute

Flatten the weights

1a	1b	1c
2a	2b	2c
3a	3b	3c

1a	1b	1c
2a	2b	2c
3a	3b	3c

0	1
-1	0

0	0
0	0

weight_flat

← get diag in $cin = 0$


← ignore $cin = 1$

1 $count = 0$

Flatten the weights

1a	1b	1c
2a	2b	2c
3a	3b	3c

1a	1b	1c
2a	2b	2c
3a	3b	3c





1a	1b	1c
2a	2b	2c
3a	3b	3c

0	0
0	0


weight_flat							
0	1	-1	0	0	0	0	0

Flatten the weights



mat-vec
multiply



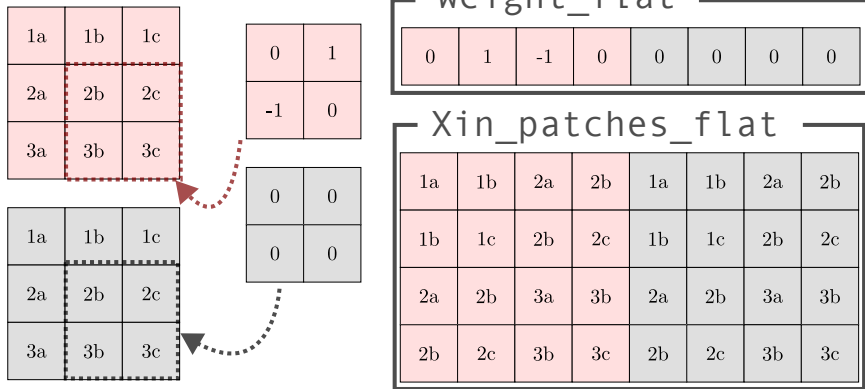
Flatten the weights



Flatten the weights

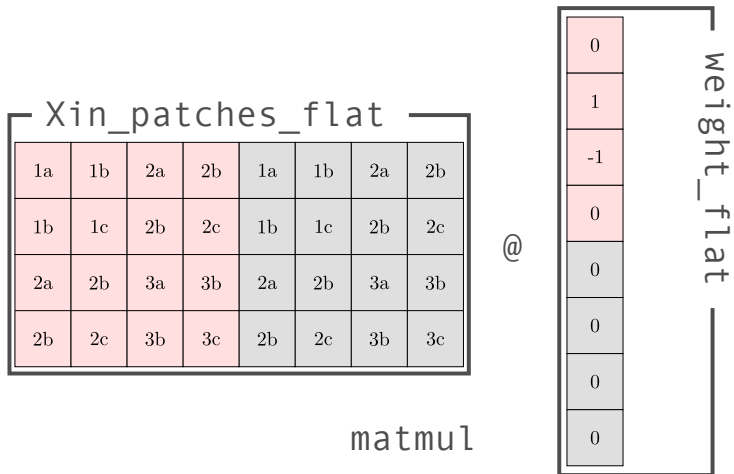


Flatten the weights



$i=0 \rightarrow 2 \times 3 \times 3 \rightarrow 4 \times 8$ $2 \cdot 2 \cdot 2$
num-patches C H W size
kernel

Sliding window via matmul



Sliding window via matmul

in ch0 diag finder



Xin_patches_flat

1a	1b	2a	2b	1a	1b	2a	2b
1b	1c	2b	2c	1b	1c	2b	2c
2a	2b	3a	3b	2a	2b	3a	3b
2b	2c	3b	3c	2b	2c	3b	3c

@

weight_flat

0	0
1	0
-1	0
0	0
0	1
0	0
0	-1
0	0

Cout

horizontal
edge
finder
in
Ch 1

output
channel

matmul

Output

$(N, P, Cout)$
↖ $H_{out} * W_{out}$

Sliding window via matmul

A filter can have both be larger than 1 many-to-many

C_{in} C_{out}

Xin_patches_flat							
1a	1b	2a	2b	1a	1b	2a	2b
1b	1c	2b	2c	1b	1c	2b	2c
2a	2b	3a	3b	2a	2b	3a	3b
2b	2c	3b	3c	2b	2c	3b	3c


@

0	0	weight_flat
1	0	
-1	0	
0	0	
0	1	
0	0	
0	-1	
0	0	
Cout		


how to get this?

matmul

Flattening Xin into Xin_flat



Flattening Xin into Xin_flat



Flattening Xin into Xin_flat

Recall patch_idx=0:

1a	1b	1c
2a	2b	2c
3a	3b	3c

1a	1b	1c
2a	2b	2c
3a	3b	3c



Our goal:

Xin_patches_flat							
1a	1b	2a	2b	1a	1b	2a	2b


What we have:

Xin_flat																	
1a	1b	1c	2a	2b	2c	3a	3b	3c	1a	1b	1c	2a	2b	2c	3a	3b	3c


The “im2col” matrix




The “im2col” matrix



The “im2col” matrix




The “im2col” matrix




The “im2col” matrix


Xin_flat @



Our goal:



What we have:



The “im2col” matrix

Recall `patch_idx=1`:

1a	1b	1c
2a	2b	2c
3a	3b	3c

1a	1b	1c
2a	2b	2c
3a	3b	3c

Our goal:

`Xin_patches_flat`

1b	1c	2b	2c	1b	1c	2b	2c
----	----	----	----	----	----	----	----


What we have:

`Xin_flat`

1a	1b	1c	2a	2b	2c	3a	3b	3c	1a	1b	1c	2a	2b	2c	3a	3b	3c
----	----	----	---------------	----	----	----	----	----	----	----	----	----	----	----	----	----	----

The “im2col” matrix

Xin_flat @



Our goal:

Xin_patches_flat

1b	1c	2b	2c	1b	1c	2b	2c
----	----	----	----	----	----	----	----

What we have:


Xin_flat

1a	1b	1c	2a	2b	2c	3a	3b	3c	1a	1b	1c	2a	2b	2c	3a	3b	3c
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----



The “im2col” matrix

Xin_flat @



Our goal:

Xin_patches_flat

2a	2b	3a	3b	2a	2b	3a	3b
----	----	----	----	----	----	----	----


What we have:

Xin_flat

1a	1b	1c	2a	2b	2c	3a	3b	3c	1a	1b	1c	2a	2b	2c	3a	3b	3c
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----



The “im2col” matrix




=

Xin_im2col

1a	1b	2a	2b	1a	1b	2a	2b
----	----	----	----	----	----	----	----

The “im2col” matrix




=


`Xin_im2col`

1a	1b	2a	2b	1a	1b	2a	2b	1b	1c	2b	2c	1b	1c	2b	2c
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----


The “im2col” matrix



The “im2col” matrix



The “im2col” matrix



The “im2col” matrix


Need to figure out where to put 1's

Xin_flat @



im2col

=

Xin_im2col



patch_position




0 for pid = 0

output_index =

$\text{patch_idx} * \text{patch_size} + \text{patch_position}$


$C_{in} \cdot K \cdot K$
11
8








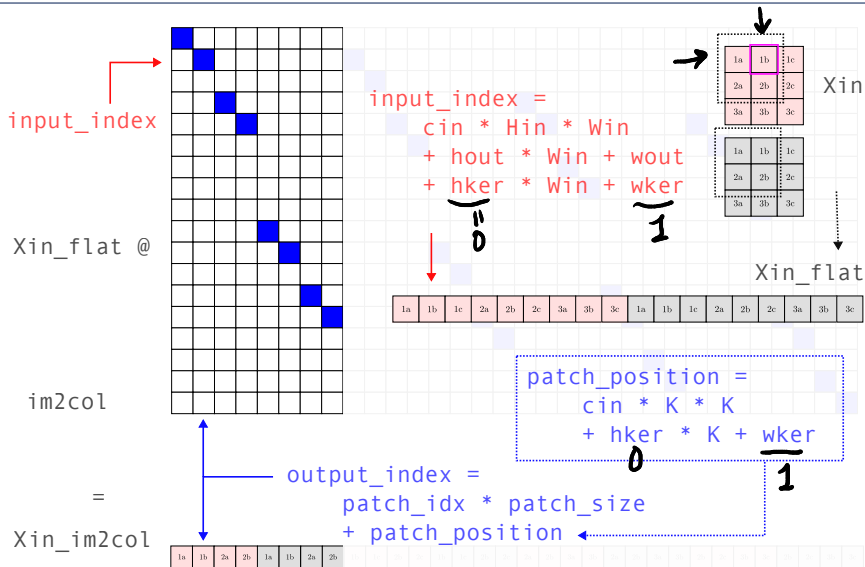
The “im2col” matrix




The “im2col” matrix




The “im2col” matrix




The “im2col” matrix




The “im2col” matrix




The “im2col” matrix



The “im2col” matrix



The “im2col” matrix



Putting it all together

stride
create the
im2col

```
1 def im2col_matrix_dense(Xin, K, S=1):
2     N, Cin, Hin, Win = Xin.shape
3     Hout, Wout = Hin - K + 1, Win - K + 1
4     P = Hout * Wout # Total number of patches per image
5     patch_size = Cin * K * K # Size of each flattened patch
6     im2col_mat_dense = np.zeros((Cin * Hin * Win, P * patch_size))
7
8     # [main loop on next slide...]
9
10    return im2col_mat_dense
```

Putting it all together

```
1  # [continued from previous slide...]
2  patch_idx = 0
3  for hout in range(Hout):
4      for wout in range(Wout):
5          for cin in range(Cin):
6              for hker in range(K):
7                  for wker in range(K):
8                      input_index = cin * Hin * Win + hout * Win +
wout + hker * Win + wker
9                      patch_position = cin * K * K + hker * K + wker
10                     output_index = patch_idx * patch_size +
patch_position
11                     im2col_mat_dense[input_index, output_index] = 1
12                     patch_idx += 1
```

which row
to place a 1

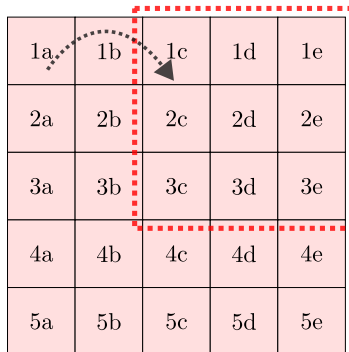
which column to place the 1

Stride = 2

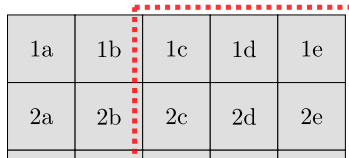
1a	1b	1c	1d	1e
2a	2b	2c	2d	2e
3a	3b	3c	3d	3e
4a	4b	4c	4d	4e
5a	5b	5c	5d	5e

1a	1b	1c	1d	1e
2a	2b	2c	2d	2e

Stride = 2





1a	1b	1c	1d	1e
2a	2b	2c	2d	2e
3a	3b	3c	3d	3e
4a	4b	4c	4d	4e
5a	5b	5c	5d	5e



1a	1b	1c	1d	1e
2a	2b	2c	2d	2e

Stride = 2



Implement im2col_matrix_dense with stride

See `lec09-in-class-ex1-im2col.ipynb`

End of part 1.

Hint 1:

$$H_{out} = \frac{H_{in} - K}{S} + 1$$

Hint 2:

w_{out}

h_{out}

multiplied

by something

Same for

$$W_{out} = \frac{W_{in} - K}{S} + 1$$

Goal for today: implement LeNet5





Image from LeCun et al. 1998

$$\text{im2col.shape} = C_{\text{in}} \cdot H_{\text{in}} \cdot W_{\text{in-by-P}} \cdot \cancel{C_{\text{out}}} \cdot K^2$$

dense
versus
sparse

$X_{\text{in_flat}}$ @
im2col
=

$X_{\text{in_im2col}}$



← mostly zero

only a
single
1 in
every
col

im2col

1a	1b	2a	2b	1a	1b	2a	2b	1b	1c	2b	2c	1b	1c	2b	2c	2a	2b	3a	3b	2a	2b	3a	3b	2b	3c	2b	2c	3b	3c
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

$$\text{im2col.shape} = C_{\text{in}} \cdot H_{\text{in}} \cdot W_{\text{in}}\text{-by-}P \cdot C_{\text{in}} \cdot K^2$$

For 1st convolution layer in LeNet5.

- $H_{\text{in}} = W_{\text{in}} = 32$
- $C_{\text{in}} = 1$
- $K = 5$
- $S = 1$
- $H_{\text{out}} = H_{\text{in}} - K + 1 = 28$
- $W_{\text{out}} = W_{\text{in}} - K + 1 = 28$
- $P = 28 * 28 \times 25$
- Final shape of im2col = 1024 -by- 19600
- 20,070,400 float64's to store ≈ 160 MB

$$(32 - 5 + 1) = 28$$

$= H_{\text{out}}$
 $= W_{\text{out}}$

not very
scalable

Sparse matrices

Pros


everything else assume to zero

- Store just the non-zero values of the matrix (memory efficient)
- Faster than dense matmul when matrices are mostly zeros (speedier)

Cons

- Irregular memory access (less cache-optimization friendly)
- Nightmare for GPUs (as of now, to the best of my knowledge)

from scipy.sparse import csr_matrix



```
1 data = [1,1,1,1,1,...] # len(data) = P * patch_size
2 row_indices = [0,1,3,4,9,...]
3 col_indices = [0,1,2,3,4,...]
4 im2col_mat_sparse = csr_matrix((data,
5                                (row_indices, col_indices)),
6                                shape=(n_rows, n_cols))
```


Example

```
1 data = [1, 2, 3, 1]
2 row_indices = [1, 2, 0, 1]
3 col_indices = [0, 1, 2, 1]
4 sparse_mat_example = csr_matrix((data, (row_indices, col_indices)),
    shape=(3, 3))
5 sparse_mat_example.toarray()
```

Output:

↓

```
1 array([[0, 0, 3],
2        [1, 1, 0],
3        [0, 2, 0]])
```

converts back to numpy

Dense-sparse multiplication

```
1 X = np.arange(9).reshape(3,3)
2 ## this won't work
3 # np.matmul(X, sparse_mat_example)
4 ## this works bc method resolution order
5 X @ sparse_mat_example
```

toarray()



```
1 array([[0, 1, 2],      array([[0, 0, 3],      array([[ 1,  5,  0],
2       [3, 4, 5],      @       [1, 1, 0],      =       [ 4, 14,  9],
3       [6, 7, 8]])      [0, 2, 0]])      [ 7, 23, 18]])
```

Implement `im2col_matrix_sparse`

- `lec09-in-class-ex2-sparse.ipynb`

modify `im2col_matrix_dense`

$$\underbrace{X_{in_flat}}_{\text{dense}} @ \text{im2col}^{\text{mat}} = X_{in_patches_flat}$$

Quick update

Can add/mul constants both on left and right now.

```
1     def __radd__(self, other):
2         return self + other
3
4     def __rmul__(self, other):
5         return self * other
6
7     def __rsub__(self, other):
8         return (-self) + other
9
10    def __rtruediv__(self, other):
11        return ag.Tensor(other) / self
```

Quick update

$\text{np.moveaxis}(X, s, d)$

if $s = -1$ $d = -2$
 $\equiv \text{transpose}$

```
1 def moveaxis(input, source, destination):
2     output = ag.Tensor(np.moveaxis(input.value, source, destination)
3     , inputs=[input], op="moveaxis")
4
5     def _backward():
6         input.grad += np.moveaxis(output.grad, source, destination)
7         return None
8     output._backward = _backward
9     return output
```

29. moveaxis

Enabling autograd for dense-sparse matmul

- Problem: our current autograd-enabled `matmul` only supports when both inputs are numpy arrays
- Observation: we don't need to track gradients on the `im2col_mat`
- Solution:

does not
change

in
training

X $X_{in} @ im2col$
✓ $sparsematmul(X_{in}, im2col)$

```
1  # [sp]arse [c]onstant (non-AG-enabled) [mat]rix [mul]tiplication
2  def spmatmul(input, sparse_mat: csr_matrix):
3      output = ag.Tensor(input.value @ sparse_mat,
4                          inputs = [input],
5                          op="spmatmul")
6
7      def _backward():
8          input.grad += output.grad @ sparse_mat.T
9          return None
10
11     output._backward = _backward
12     return output
```

don't update
sparse_mat
grad


Exercise 3

lec 09_data.csv




tiny-mnist

- Create the Conv2d layer ←
- Create the AvgPool2d layer ←
- Create the LeNet5 model ←



- Create the `nn.CrossEntropyLoss` ←

Sanity check for conv2d



hints:


move axis

spcmatmul


from EX1:

Sanity check for pool

Cont = 0 is the average of
 $C_{in} = 0$
and $Cont = 1$ $C_{in} = 1$
and so on.
 $K = S$



$C_{in} = 0$
I just average
the value
in patch
in C_{in} to $Cont$



like a conv 2d
except

$$C_{in} = Cont$$

Sanity check for training run

You should achieve 100 percentage training accuracy

References I

- [LeC+98] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. “Gradient-based learning applied to document recognition”. In: *Proceedings of the IEEE* 86.11 (1998), pp. 2278–2324.