

Deep Learning in Practice

Lecture 10 — CS 577 Deep Learning

Instructor: Yutong Wang

Computer Science
Illinois Institute of Technology

October 23, 2024

The gap between ag.Tensor and PyTorch

- `ag.Tensor` assume everything has gradients
- many layers/tools are not yet supported:
 - advanced optimizers
 - data wrangling tools like mini-batching
 - normalization (e.g., batch norm)
 - dropout
 - weight decay
 - model compression
 - ...
- and many other shortcoming

Grads

```
1  class Tensor: # Tensor with grads
2      def __init__(self,
3                    value,
4                    op="",
5                    _backward= lambda : None,
6                    inputs=[],
7                    label=""):
8
9      if type(value) in [float ,int]:
10         value = np.array(value)
11         self.value = value
12
13         self.grad = np.zeros_like(value) # <- always need this??
```

PyTorch tensors has

```
1  torch.Tensor.requires_grad_ # true if needs to be updated
```

The gap between ag.Tensor and PyTorch

- ag.Tensor assume everything has gradients
- many layers/tools are not yet supported:
 - **advanced optimizers**
 - data wrangling tools like mini-batching
 - normalization (e.g., batch norm)
 - dropout
 - weight decay
 - model compression
 - ...
- and many other shortcoming

Stochastic gradient descent (SGD)

Let $\eta_t > 0$ be learning rates, $t = 1, 2, \dots$

Let $m \geq 1$ be an integer

- Initialize $\boldsymbol{\theta}$
- While not converged ($t = \text{iteration counter}$):
 - Select m samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ and matching labels $\{y^{(1)}, \dots, y^{(m)}\}$
 - Compute gradient $\mathbf{g} \leftarrow \nabla_{\boldsymbol{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}, \boldsymbol{\theta}), y^{(i)})$
 - Compute update $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \eta_t \mathbf{g}$

Stochastic gradient descent (SGD)

Let $\eta_t > 0$ be learning rates, $t = 1, 2, \dots$

Let $m \geq 1$ be an integer

- Initialize $\boldsymbol{\theta}^{(0)}$
- While not converged ($t =$ iteration counter):
 - Select m samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ and matching labels $\{y^{(1)}, \dots, y^{(m)}\}$
 - Compute gradient $\mathbf{g}^{(t)} \leftarrow \nabla_{\boldsymbol{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}, \boldsymbol{\theta}^{(t-1)}), y^{(i)})$
 - Compute update $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta_t \mathbf{g}^{(t)}$

“Unrolling” SGD

$$\boldsymbol{\theta}^{(1)} \leftarrow \boldsymbol{\theta}^{(0)} - \eta_1 \mathbf{g}^{(1)}$$

“Unrolling” SGD

$$\boldsymbol{\theta}^{(2)} \leftarrow \boldsymbol{\theta}^{(1)} - \eta_2 \mathbf{g}^{(2)} = \boldsymbol{\theta}^{(0)} - \eta_1 \mathbf{g}^{(1)} - \eta_2 \mathbf{g}^{(2)}$$

“Unrolling” gradient descent

$$\boldsymbol{\theta}^{(3)} \leftarrow \boldsymbol{\theta}^{(2)} - \eta_1 \mathbf{g}^{(3)} = \boldsymbol{\theta}^{(0)} - \eta_1 \mathbf{g}^{(1)} - \eta_2 \mathbf{g}^{(2)} - \eta_3 \mathbf{g}^{(3)}$$

“Unrolling” gradient descent

$$\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta_t \mathbf{g}^{(t)} = \boldsymbol{\theta}^{(0)} - \eta_1 \mathbf{g}^{(1)} - \eta_2 \mathbf{g}^{(2)} - \eta_3 \mathbf{g}^{(3)} \dots - \eta_t \mathbf{g}^{(t)}$$

SGD with Momentum

Let $\eta_t > 0$ and $t = 1, 2, \dots$ be the iteration counter. Let $\gamma \in [0, 1)$ be the momentum parameter.

- Initialize $\boldsymbol{\theta}^{(0)}$ and $\boxed{\mathbf{v}^{(0)} = \mathbf{0}}$ (momentum vector)
- While not converged ($t = \text{iteration counter}$):
 - Select m samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ and matching labels $\{y^{(1)}, \dots, y^{(m)}\}$
 - Compute gradient $\mathbf{g}^{(t)} \leftarrow \nabla_{\boldsymbol{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}), \boldsymbol{\theta}^{(t-1)}), y^{(i)})$
 - Update velocity $\boxed{\mathbf{v}^{(t)} \leftarrow \gamma \mathbf{v}^{(t-1)} + \mathbf{g}^{(t)}}$
 - Update parameters $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta_t \mathbf{v}^{(t)}$

[Sut+13] “On the importance of initialization and momentum in deep learning”

“Unrolling” Velocity in SGD with Momentum

SGD with momentum: just the key components

- Initialize $\boldsymbol{\theta}^{(0)}$ and $\boxed{\mathbf{v}^{(0)} = 0}$ (momentum vector)

- Update velocity $\boxed{\mathbf{v}^{(t)} \leftarrow \gamma \mathbf{v}^{(t-1)} + \mathbf{g}^{(t)}}$

- Update parameters $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta_t \mathbf{v}^{(t)}$

$$\mathbf{v}^{(1)} \leftarrow \mathbf{g}^{(1)}$$

“Unrolling” Velocity in SGD with Momentum

SGD with momentum: just the key components

- Initialize $\theta^{(0)}$ and $\mathbf{v}^{(0)} = 0$ (momentum vector)
- Update velocity $\mathbf{v}^{(t)} \leftarrow \gamma \mathbf{v}^{(t-1)} + \mathbf{g}^{(t)}$
- Update parameters $\theta^{(t)} \leftarrow \theta^{(t-1)} - \eta_t \mathbf{v}^{(t)}$

$$\mathbf{v}^{(1)} \leftarrow \mathbf{g}^{(1)}$$

$$\mathbf{v}^{(2)} \leftarrow \gamma \mathbf{v}^{(1)} + \mathbf{g}^{(2)} = \gamma \mathbf{g}^{(1)} + \mathbf{g}^{(2)}$$

“Unrolling” Velocity in SGD with Momentum

SGD with momentum: just the key components

- Initialize $\theta^{(0)}$ and $\boxed{\mathbf{v}^{(0)} = 0}$ (momentum vector)

- Update velocity $\boxed{\mathbf{v}^{(t)} \leftarrow \gamma \mathbf{v}^{(t-1)} + \mathbf{g}^{(t)}}$

- Update parameters $\theta^{(t)} \leftarrow \theta^{(t-1)} - \eta_t \mathbf{v}^{(t)}$

$$\mathbf{v}^{(1)} \leftarrow \mathbf{g}^{(1)}$$

$$\mathbf{v}^{(2)} \leftarrow \gamma \mathbf{v}^{(1)} + \mathbf{g}^{(2)} = \gamma \mathbf{g}^{(1)} + \mathbf{g}^{(2)}$$

$$\mathbf{v}^{(3)} \leftarrow \gamma \mathbf{v}^{(2)} + \mathbf{g}^{(3)} = \gamma^2 \mathbf{g}^{(1)} + \gamma \mathbf{g}^{(2)} + \mathbf{g}^{(3)}$$

“Unrolling” Velocity in SGD with Momentum

SGD with momentum: just the key components

- Initialize $\theta^{(0)}$ and $\boxed{\mathbf{v}^{(0)} = 0}$ (momentum vector)
- Update velocity $\boxed{\mathbf{v}^{(t)} \leftarrow \gamma \mathbf{v}^{(t-1)} + \mathbf{g}^{(t)}}$
- Update parameters $\theta^{(t)} \leftarrow \theta^{(t-1)} - \eta_t \mathbf{v}^{(t)}$

$$\mathbf{v}^{(1)} \leftarrow \mathbf{g}^{(1)}$$

$$\mathbf{v}^{(2)} \leftarrow \gamma \mathbf{v}^{(1)} + \mathbf{g}^{(2)} = \gamma \mathbf{g}^{(1)} + \mathbf{g}^{(2)}$$

$$\mathbf{v}^{(3)} \leftarrow \gamma \mathbf{v}^{(2)} + \mathbf{g}^{(3)} = \gamma^2 \mathbf{g}^{(1)} + \gamma \mathbf{g}^{(2)} + \mathbf{g}^{(3)}$$

$$\mathbf{v}^{(t)} \leftarrow \gamma \mathbf{v}^{(t-1)} + \mathbf{g}^{(t)} = \gamma^{t-1} \mathbf{g}^{(1)} + \gamma^{t-2} \mathbf{g}^{(2)} + \dots + \mathbf{g}^{(t)}$$


SGD with Momentum

Let $\eta_t > 0$ and $t = 1, 2, \dots$ be the iteration counter. Let $\gamma \in [0, 1)$ be the momentum parameter.

- Initialize $\boldsymbol{\theta}^{(0)}$ and $\boxed{\mathbf{v}^{(0)} = 0}$ (momentum)
- While not converged ($t = \text{iteration counter}$):
 - Select m samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ and matching labels $\{y^{(1)}, \dots, y^{(m)}\}$
 - Compute gradient $\mathbf{g}^{(t)} \leftarrow \nabla_{\boldsymbol{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}), \boldsymbol{\theta}^{(t-1)}), y^{(i)})$
 - Update velocity $\boxed{\mathbf{v}^{(t)} \leftarrow \gamma \mathbf{v}^{(t-1)} + \mathbf{g}^{(t)}}$
 - Update parameters $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta_t \mathbf{v}^{(t)}$

In class exercise

lec10-in-class-ex-optim.ipynb



Adam Optimizer

Let $\eta_t > 0$, $\beta_1, \beta_2 \in [0, 1)$, and $t = 1, 2, \dots$ be the iteration counter.

- Initialize $\boldsymbol{\theta}^{(0)}$, $\boxed{\mathbf{v}^{(0)} = 0}$ (first moment), $\boxed{\mathbf{s}^{(0)} = 0}$ (second moment)
- While not converged ($t =$ iteration counter):
 - Select m samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ and matching labels $\{y^{(1)}, \dots, y^{(m)}\}$
 - Compute gradient $\mathbf{g}^{(t)} \leftarrow \nabla_{\boldsymbol{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}, \boldsymbol{\theta}^{(t-1)}), y^{(i)})$
 - Update first moment: $\boxed{\mathbf{v}^{(t)} \leftarrow \beta_1 \mathbf{v}^{(t-1)} + (1 - \beta_1) \mathbf{g}^{(t)}}$ $\boxed{\hat{\mathbf{v}}^{(t)} \leftarrow \frac{\mathbf{v}^{(t)}}{1 - \beta_1^t}}$
 - Update second moment: $\boxed{\mathbf{s}^{(t)} \leftarrow \beta_2 \mathbf{s}^{(t-1)} + (1 - \beta_2) (\mathbf{g}^{(t)})^2}$ $\boxed{\hat{\mathbf{s}}^{(t)} \leftarrow \frac{\mathbf{s}^{(t)}}{1 - \beta_2^t}}$
 - Update parameters: $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta_t \hat{\mathbf{v}}^{(t)} / (\sqrt{\hat{\mathbf{s}}^{(t)}} + \epsilon)$

The gap between ag.Tensor and PyTorch

- ag.Tensor assume everything has gradients
- many layers/tools are not yet supported:
 - advanced optimizers
 - **data wrangling tools like mini-batching**
 - normalization (e.g., batch norm)
 - dropout
 - weight decay
 - model compression
 - ...
- and many other shortcoming

Data minibatching

From lenet5_mnist.ipynb

```
1 transform = transforms.Compose([transforms.ToTensor(), transforms.  
    Normalize((0.5,), (0.5,))])  
2  
3 mnist_train = torchvision.datasets.MNIST(root='./data', train=True,  
    download=True, transform=transform)  
4  
5 mnist_train_loader = DataLoader(mnist_train, batch_size=64, shuffle=True  
    )
```

Stochastic gradient descent (SGD)

Let $\eta_t > 0$ be learning rates, $t = 1, 2, \dots$

Let $m \geq 1$ be an integer

- Initialize $\boldsymbol{\theta}$
- While not converged ($t = \text{iteration counter}$):
 - Select m samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ and matching labels $\{y^{(1)}, \dots, y^{(m)}\}$
 - Compute gradient $\mathbf{g} \leftarrow \nabla_{\boldsymbol{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}, \boldsymbol{\theta}), y^{(i)})$
 - Compute update $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} - \eta_t \mathbf{g}$

Stochastic gradient descent (SGD)

Let $\eta_t > 0$ be learning rates, $t = 1, 2, \dots$

Let $m \geq 1$ be an integer

- Initialize $\boldsymbol{\theta}^{(0)}$
- While not converged ($t =$ iteration counter):
 - Select m samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ and matching labels $\{y^{(1)}, \dots, y^{(m)}\}$
 - Compute gradient $\mathbf{g}^{(t)} \leftarrow \nabla_{\boldsymbol{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}, \boldsymbol{\theta}^{(t-1)}), y^{(i)})$
 - Compute update $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta_t \mathbf{g}^{(t)}$

Stochastic gradient descent (SGD)

Let $\eta_t > 0$ be learning rates, $t = 1, 2, \dots$

Let $m \geq 1$ be an integer

- Initialize $\boldsymbol{\theta}^{(0)}$
- While not converged ($t = \text{iteration counter}$):
 - Select m samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ and matching labels $\{y^{(1)}, \dots, y^{(m)}\}$
 - Compute gradient $\mathbf{g}^{(t)} \leftarrow \nabla_{\boldsymbol{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}, \boldsymbol{\theta}^{(t-1)}), y^{(i)})$
 - Compute update $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta_t \mathbf{g}^{(t)}$

Stochastic gradient descent (SGD)

Let $\eta_t > 0$ be learning rates, $t = 1, 2, \dots$

Let $m \geq 1$ be an integer

- Initialize $\boldsymbol{\theta}^{(0)}$
- While not converged ($t =$ iteration counter):
 - Select m samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ and matching labels $\{y^{(1)}, \dots, y^{(m)}\}$
 - Compute gradient $\mathbf{g}^{(t)} \leftarrow \nabla_{\boldsymbol{\theta}} \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}, \boldsymbol{\theta}^{(t-1)}), y^{(i)})$
 - Compute update $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta_t \mathbf{g}^{(t)}$

```
1 mnist_train_loader = DataLoader(mnist_train, batch_size=64, shuffle=True)
```


Test data

```
1 mnist_test = torchvision.datasets.MNIST(root='./data', train=False,  
    download=True, transform=transform)  
2  
3 mnist_test_loader = DataLoader(mnist_test, batch_size=64, shuffle=False)
```

Don't look at the test data...

...when you are

- deciding on an architecture to use
- tuning hyperparameters
- training the data

- Because...

Test data

```
1 mnist_test = torchvision.datasets.MNIST(root='./data', train=False,  
    download=True, transform=transform)  
2  
3 mnist_test_loader = DataLoader(mnist_test, batch_size=64, shuffle=False)
```

Don't look at the test data...

...when you are

- deciding on an architecture to use
 - tuning hyperparameters
 - training the data
-
- Because...
 - ...You might trick yourself (and others) into making tweaks that are helping

Test data

```
1 mnist_test = torchvision.datasets.MNIST(root='./data', train=False,  
    download=True, transform=transform)  
2  
3 mnist_test_loader = DataLoader(mnist_test, batch_size=64, shuffle=False)
```

Don't look at the test data...

...when you are

- deciding on an architecture to use
 - tuning hyperparameters
 - training the data
-
- Because...
 - ...You might trick yourself (and others) into making tweaks that are helping
 - ...In practice you don't have the test data ahead of time

Real world example

- Goal: you're helping the US Postal Service to improve mail delivery (it's the early 1990s)



Real world example

- Goal: you're helping the US Postal Service to improve mail delivery (it's the early 1990s)



- Collect *training data* and hire people to label handwritten digits



Real world example

- Goal: you're helping the US Postal Service to improve mail delivery (it's the early 1990s)





- Collect *training data* and hire people to label handwritten digits
- Deploy in the wild (on test data that the model has never seen before)

Overfitting

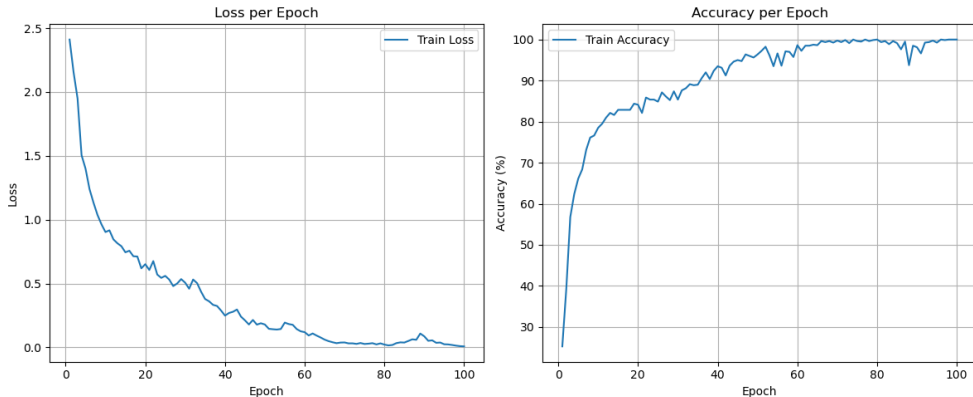


Overfitting




-
- Test accuracy: 81.1% accuracy

Overfitting



-
- Test accuracy: 81.1% accuracy
- Disclaimer: I injected noise into the training data (but not the test data) to prove a point. MNIST is very easy. Even LeNet5 can reach 98% test accuracy.

Overfitting




-
- **Test accuracy:** 81.1% accuracy
- Disclaimer: I injected noise into the training data (but not the test data) to prove a point. MNIST is very easy. Even LeNet5 can reach 98% test accuracy.
- Solution: split the training data further into train and validation sets.


Train/valid splitting

```
1 import torch.utils.data as data
2 # [...]
3 train_size = int(0.8 * len(train_dataset))
4 val_size = len(train_dataset) - train_size
5 train_dataset, val_dataset = data.random_split(train_dataset, [
    train_size, val_size])
```

Split train into train/validation sets




Split train into train/validation sets

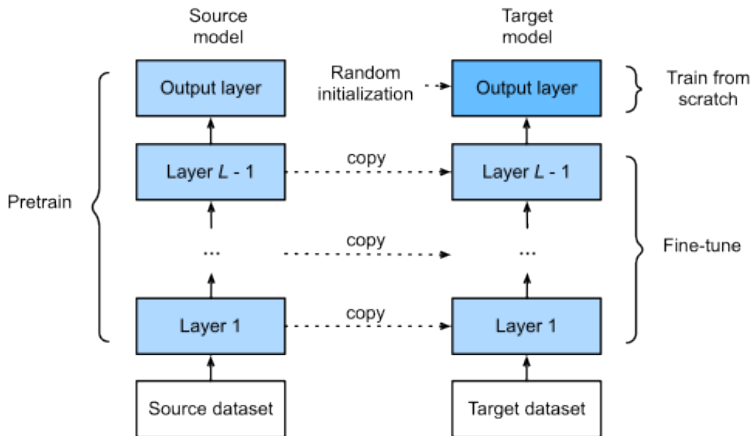


- 85.8% test accuracy at epoch with the highest validation accuracy.

A data challenge for today



Transfer learning



From <https://d2l.ai/> chapter on fine-tuning

Transfer learning

- Train model on one task
- Copy the weights to a new task
- Update weights based on optimization from new task

```
1 class LeNet5_Custom(nn.Module):  
2     def __init__(self, pretrained_model):  
3         super(LeNet5_Custom, self).__init__()  
4         # Clone the layers from the pre-trained model  
5         self.conv1 = nn.Conv2d(1, 6, kernel_size=5)  
6         self.conv1.weight = nn.Parameter(pretrained_model.conv1.weight.  
clone())  
7         self.conv1.bias = nn.Parameter(pretrained_model.conv1.bias.clone  
(()))
```


Other ideas

- Data augmentation
- Batch normalization
- Weight decay
- Skip connections
- Dropout

References I

- [Sut+13] Ilya Sutskever, James Martens, George Dahl, and Geoffrey Hinton. “On the importance of initialization and momentum in deep learning”. In: *International conference on machine learning*. PMLR. 2013, pp. 1139–1147.