

Sequence models

Lecture 11 — CS 577 Deep Learning

Instructor: Yutong Wang

Computer Science
Illinois Institute of Technology

October 30, 2024

- Seq model
- Flash Attention

Systems (ML + optimizing for memory bandwidth)
memory-bound

moving things

from
memory

to compute
(GPU)

is the
bottleneck

- Andre Bauer /
time series
for AI

Minxuan Zhou
Hardware topics

Tailored hardware
In-memory compute

- { Matt Smith
CV for ecology
review for final

Sequence models

Lecture 11 — CS 577 Deep Learning

Instructor: Yutong Wang

Computer Science
Illinois Institute of Technology

October 30, 2024

understanding insects biodiversity


Topics

- Temporal convolutional networks
- Recurrent neural networks
- Transformers

Sequences

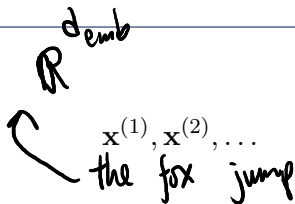
dependency across time
 $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots$

- Example: Time series (acoustic signal, weather, stock price,...)



Sequences

$\mathbb{R}^{d_{\text{emb}}}$
 $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots$
the fox jump



- Example: Time series (acoustic signal, weather, stock price,...)
- Example: Sequence of word/token embeddings

Sequences

$$\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots$$

- Example: Time series (acoustic signal, weather, stock price,...)
- Example: Sequence of word/token embeddings
- Task: denoising (make a signal less noisy)

clean up audio

Sequences

$$\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots$$

- Example: Time series (acoustic signal, weather, stock price,...)
- Example: Sequence of word/token embeddings
- Task: denoising (make a signal less noisy)
- Task: forecasting (what is the weather like next week?)

Sequences

$$\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots$$

- Example: Time series (acoustic signal, weather, stock price,...)
- Example: Sequence of word/token embeddings
- Task: denoising (make a signal less noisy)
- Task: forecasting (what is the weather like next week?)
- Task: classification (is this news article about sport or about technology?)

document
classif.



↳ generative
next word
prediction

Sequences

$$\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots$$

- Example: Time series (acoustic signal, weather, stock price,...)
- Example: Sequence of word/token embeddings
- Task: denoising (make a signal less noisy)
- Task: forecasting (what is the weather like next week?)
- Task: classification (is this news article about sport or about technology?)
- Task: regression (how favorable is this product review?)

sentiment analysis 0 → 100


Causal models

put "causal
rail guard"
in training loop
implicitly make the
model causal

$$f(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(t)}; \theta) = \mathbf{y}^{(t)}$$

Causality assumption

The output $\mathbf{y}^{(t)}$ does not depend on future values of $\mathbf{x}^{(t')}$ where $t' > t$



Prediction/forecasting/generative models

input space = output space
when is the not true?

- Let $\hat{\mathbf{x}}^{(t+1)} := \mathbf{y}^{(t)}$

-

$$\underbrace{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(C)}}_{\text{observed}} \xrightarrow{f} \underbrace{\hat{\mathbf{x}}^{(C+1)}, \hat{\mathbf{x}}^{(C+2)}, \dots}_{\text{predicted/generated}}$$

- Next: for simplicity, let's consider $x^{(t)} \in \mathbb{R}$. (So drop the bold fontface.)

Context = ↑ time observed
max up to

Prediction/forecasting/generative models

- Let $\hat{\mathbf{x}}^{(t+1)} := \mathbf{y}^{(t)}$

-

$$\underbrace{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(C)}}_{\text{observed}} \xrightarrow{f} \underbrace{\hat{\mathbf{x}}^{(C+1)}, \hat{\mathbf{x}}^{(C+2)}, \dots}_{\text{predicted/generated}}$$

- Next: for simplicity, let's consider $x^{(t)} \in \mathbb{R}$. (So drop the bold fontface.)

Scalar

simplest non-trivial

Later: how this leads to RNNs

Autoregressive model

AR(5)

- For $t \geq s$ ← $s := \text{filter-size}$

$\hat{x}^{(t+1)}$

$$= y^{(t)} = f(x^{(1)}, \dots, x^{(t)}; w_0, \dots, w_{s-1}, b)$$

$s=5$

Goal

$$= \sum_{\tau=0}^{s-1} w_{\tau} x^{(t-\tau)} + b$$

$x^{(1)}, \dots,$



weight / slope



Autoregressive model

- For $t \geq s$

$$\hat{x}^{(t+1)} = y^{(t)} = f(x^{(1)}, \dots, x^{(t)}; w_0, \dots, w_{s-1}, b)$$
$$= \sum_{\tau=0}^{s-1} w_{\tau} x^{(t-\tau)} + b$$

- Loss

$$\frac{1}{C} \sum_{t=1}^C (x^{(t+1)} - \underbrace{f(x^{(1)}, \dots, x^{(t)}; w_0, \dots, w_{s-1}, b)}_{\text{predict / forecast}})^2$$

next time data arrives

↙ MSE

Autoregressive model

- For $t \geq s$

$$\begin{aligned}y^{(t)} &= f(x^{(1)}, \dots, x^{(t)}; w_0, \dots, w_{s-1}, b) \\ &= \sum_{\tau=0}^{s-1} w_{\tau} x^{(t-\tau)} + b\end{aligned}$$

- Loss

$$\frac{1}{C} \sum_{t=1}^C (x^{(t+1)} - f(x^{(1)}, \dots, x^{(t)}; w_0, \dots, w_{s-1}, b))^2$$

- Solve for $\{w_0, \dots, w_{s-1}, b\}$

↳ np.array

Autoregressive model

- For $t \geq s$

$$\begin{aligned}y^{(t)} &= f(x^{(1)}, \dots, x^{(t)}; w_0, \dots, w_{s-1}, b) \\ &= \sum_{\tau=0}^{s-1} w_{\tau} x^{(t-\tau)} + b\end{aligned}$$

- Loss

$$\frac{1}{C} \sum_{t=1}^C (x^{(t+1)} - f(x^{(1)}, \dots, x^{(t)}; w_0, \dots, w_{s-1}, b))^2$$

- Solve for w_0, \dots, w_{s-1}, b

- Next: seq2col

im2col

Seq2col: in class exercise 1

In class exercise

- Complete the seq2col function
- Discuss with your neighbors regarding questions in the “Fibonacci” block

Seq2col: in class exercise 1 discussion

$$x^{(t+1)} = x^{(t)} + x^{(t-1)}$$

1 1 2 3 5 8

How do you interpret the coefficients when we set `filter_size = 3`?

```
1 filter_size = 3
2 def fibonacci(n):
3     [...]
4
5 fibonnaci_seq = list(fibonacci(20))
6
7 X, y = seq2col(fibonnaci_seq[:10], filter_size)
8
9 X_tilde = np.hstack([X, np.ones((X.shape[0], 1))])
10 w = np.linalg.pinv(X_tilde) @ y
11 np.round(w,5)
12 # array([-0.,  1.,  1.,  0.])
```

least w b most recent

Another way to look at autoregressive models

- Drop the bias term for simplicity

↳ to RNN (recurrent neural network)

$$y^{(t)} = w_0 x^{(t)} + \underbrace{\sum_{\tau=1}^{s-1} w_{\tau} x^{(t-\tau)}}_{\text{weight}} = w_0 x^{(t)} + \underbrace{\begin{bmatrix} w_1 \\ \vdots \\ w_{s-1} \end{bmatrix}^{\top} \begin{bmatrix} x^{(t-1)} \\ \vdots \\ x^{(t-(s-1))} \end{bmatrix}}_{\text{dot product}}$$

Another way to look at autoregressive models

- Drop the bias term for simplicity

$$y^{(t)} = w_0 x^{(t)} + \sum_{\tau=1}^{s-1} w_{\tau} x^{(t-\tau)} = w_0 x^{(t)} + \begin{bmatrix} w_1 \\ \vdots \\ w_{s-1} \end{bmatrix}^{\top} \begin{bmatrix} x^{(t-1)} \\ \vdots \\ x^{(t-(s-1))} \end{bmatrix}$$

- History up to time $t - 1$ (aka hidden state at time $t - 1$)

$$y^{(t)} = w_0 x^{(t)} + \begin{bmatrix} w_1 \\ \vdots \\ w_{s-1} \end{bmatrix}^{\top} \tilde{\mathbf{h}}^{(t-1)}$$

Υ
current input

+ weight @ hidden-state

def !!

Another way to look at autoregressive models

- Drop the bias term for simplicity

$$y^{(t)} = w_0 x^{(t)} + \sum_{\tau=1}^{s-1} w_{\tau} x^{(t-\tau)} = w_0 x^{(t)} + \begin{bmatrix} w_1 \\ \vdots \\ w_{s-1} \end{bmatrix}^{\top} \begin{bmatrix} x^{(t-1)} \\ \vdots \\ x^{(t-(s-1))} \end{bmatrix}$$

- History up to time $t - 1$ (aka hidden state at time $t - 1$)

*Instead
of modeling
 $y^{(t)}$*

$$y^{(t)} = w_0 x^{(t)} + \begin{bmatrix} w_1 \\ \vdots \\ w_{s-1} \end{bmatrix}^{\top} \tilde{\mathbf{h}}^{(t-1)}$$


hidden state

- Next: directly model the hidden state

Recurrent unit (with linear activation)

- Recurrent unit

$$\mathbf{h}^{(t)} = f(x^{(t)}, \mathbf{h}^{(t-1)}; \boldsymbol{\theta}) = \mathbf{W}_{\text{rec}} \mathbf{h}^{(t-1)} + \mathbf{w}_{\text{in}} x^{(t)}$$


hidden state input

Recurrent unit (with linear activation)

- Recurrent unit

$$d_h \quad \mathbf{h}^{(t)} = f(x^{(t)}, \mathbf{h}^{(t-1)}; \boldsymbol{\theta}) = \mathbf{W}_{\text{rec}} \mathbf{h}^{(t-1)} + \mathbf{w}_{\text{in}} x^{(t)}$$

- $\mathbf{h}^{(t)} \in \mathbb{R}^{d_h}$ (hidden) states at time t

Recurrent unit (with linear activation)

- Recurrent unit

$$\mathbf{h}^{(t)} = f(x^{(t)}, \mathbf{h}^{(t-1)}; \boldsymbol{\theta}) = \mathbf{W}_{\text{rec}}\mathbf{h}^{(t-1)} + \mathbf{w}_{\text{in}}x^{(t)}$$

- $\mathbf{h}^{(t)} \in \mathbb{R}^{d_{\text{h}}}$ (hidden) states at time t
- d_{h} hidden dimension

Recurrent unit (with linear activation)

- Recurrent unit

$$\mathbf{h}^{(t)} = f(x^{(t)}, \mathbf{h}^{(t-1)}; \theta) = \mathbf{W}_{\text{rec}} \mathbf{h}^{(t-1)} + \mathbf{w}_{\text{in}} x^{(t)}$$

- $\mathbf{h}^{(t)} \in \mathbb{R}^{d_h}$ (hidden) states at time t
- d_h hidden dimension
- Read out $\mathbf{w}_{\text{out}} \in \mathbb{R}^{d_h}$

final layer of your model
classification "head"

$$y^{(t)} = \mathbf{w}_{\text{out}}^{\top} \mathbf{h}^{(t)}$$

regression
"head"
hidden state

Recurrent unit (with linear activation)

- Recurrent unit

$$\mathbf{h}^{(t)} = f(x^{(t)}, \mathbf{h}^{(t-1)}; \boldsymbol{\theta}) = \mathbf{W}_{\text{rec}} \mathbf{h}^{(t-1)} + \mathbf{w}_{\text{in}} x^{(t)}$$

- $\mathbf{h}^{(t)} \in \mathbb{R}^{d_h}$ (hidden) states at time t
- d_h hidden dimension
- Read out $\mathbf{w}_{\text{out}} \in \mathbb{R}^{d_h}$

$$y^{(t)} = \mathbf{w}_{\text{out}}^\top \mathbf{h}^{(t)}$$

- $\mathbf{W}_{\text{rec}} \in \mathbb{R}^{d_h \times d_h}$, $\mathbf{w}_{\text{in}} \in \mathbb{R}^{d_h}$ and $\mathbf{w}_{\text{out}} \in \mathbb{R}^{d_h}$ are parameters

Recovering the autoregressive model

- Autoregressive (AR) model (with filter size 4)

$$y^{(t)} = w_0 x^{(t)} + \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix}^\top \tilde{\mathbf{h}}^{(t-1)} \quad \text{where} \quad \tilde{\mathbf{h}}^{(t-1)} = \begin{bmatrix} x^{(t-1)} \\ x^{(t-2)} \\ x^{(t-3)} \end{bmatrix}$$

- Recurrent unit: how can we choose the parameters to recover the AR model?

Want this \rightarrow

$$\mathbf{h}^{(t)} :=$$

$$\begin{bmatrix} x^{(t)} \\ x^{(t-1)} \\ x^{(t-2)} \\ x^{(t-3)} \end{bmatrix}$$

$$\mathbf{h}^{(t)} = f(x^{(t)}, \mathbf{h}^{(t-1)}; \theta) = \mathbf{W}_{\text{rec}} \mathbf{h}^{(t-1)} + \mathbf{w}_{\text{in}} x^{(t)}$$

$$\downarrow$$
$$\begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Recovering the autoregressive model

- Autoregressive (AR) model (with filter size 4)

$$y^{(t)} = w_0 x^{(t)} + \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix}^\top \tilde{\mathbf{h}}^{(t-1)} \quad \text{where} \quad \tilde{\mathbf{h}}^{(t-1)} = \begin{bmatrix} x^{(t-1)} \\ x^{(t-2)} \\ x^{(t-3)} \end{bmatrix}$$

- Recurrent unit: how can we choose the parameters to recover the AR model?

Want

$$\mathbf{h}^{(t)} = \begin{bmatrix} x^{(t)} \\ x^{(t-1)} \\ x^{(t-2)} \\ x^{(t-3)} \end{bmatrix}$$

Assume this is achieved up to time $t-1$

$$\mathbf{h}^{(t)} = f(x^{(t)}, \mathbf{h}^{(t-1)}; \theta) = \mathbf{W}_{\text{rec}} \mathbf{h}^{(t-1)} + \mathbf{w}_{\text{in}} x^{(t)}$$

then

$$\begin{bmatrix} 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \parallel \begin{bmatrix} x^{(t-1)} \\ x^{(t-2)} \\ x^{(t-3)} \\ \cancel{x^{(t-4)}} \end{bmatrix}$$

drop this

Recovering the autoregressive model

- Autoregressive (AR) model (with filter size 4)

$$y^{(t)} = \underline{w_0}x^{(t)} + \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix}^\top \tilde{\mathbf{h}}^{(t-1)} \quad \text{where} \quad \tilde{\mathbf{h}}^{(t-1)} = \begin{bmatrix} x^{(t-1)} \\ x^{(t-2)} \\ x^{(t-3)} \end{bmatrix}$$

- Recurrent unit: how can we choose the parameters to recover the AR model?

$$\mathbf{h}^{(t)} = f(x^{(t)}, \mathbf{h}^{(t-1)}; \boldsymbol{\theta}) = \mathbf{W}_{\text{rec}} \mathbf{h}^{(t-1)} + \mathbf{w}_{\text{in}} x^{(t)}$$

$$\mathbf{W}_{\text{out}} = \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_3 \end{bmatrix}$$

$$\mathbf{W}_{\text{out}}^\top \mathbf{h}^{(t)}$$

recovers
auto reg

$$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

choose

$$\begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$\xrightarrow{\text{upshot}} \mathbf{h}^{(t)} = \begin{bmatrix} x^{(t)} \\ x^{(t-1)} \\ x^{(t-2)} \\ x^{(t-3)} \end{bmatrix}$$

ensure

Recovering the autoregressive model

a single recurrent unit is already expressive

- Autoregressive (AR) model (with filter size 4)

$$y^{(t)} = w_0 x^{(t)} + \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix}^\top \tilde{\mathbf{h}}^{(t-1)} \quad \text{where} \quad \tilde{\mathbf{h}}^{(t-1)} = \begin{bmatrix} x^{(t-1)} \\ x^{(t-2)} \\ x^{(t-3)} \end{bmatrix}$$

- Recurrent unit: how can we choose the parameters to recover the AR model?

$$\mathbf{h}^{(t)} = f(x^{(t)}, \mathbf{h}^{(t-1)}; \theta) = \mathbf{W}_{\text{rec}} \mathbf{h}^{(t-1)} + \mathbf{w}_{\text{in}} x^{(t)}$$

AR(4) can be implemented as a
Rec Unit with $d_n = 4$

Recurrent unit (with linear activation)

- Recurrent unit

$$\mathbf{h}^{(t)} = f(\mathbf{x}^{(t)}, \mathbf{h}^{(t-1)}; \boldsymbol{\theta}) = \mathbf{W}_{\text{rec}} \mathbf{h}^{(t-1)} + \mathbf{W}_{\text{in}} \mathbf{x}^{(t)}$$

- \mathbf{W}_{rec} and \mathbf{W}_{in} are parameters
- $\mathbf{h}^{(t)}$ (hidden) states at time t

can be \top vectorial

Recurrent unit (with linear activation)


- Recurrent unit

$$\mathbf{h}^{(t)} = f(\mathbf{x}^{(t)}, \mathbf{h}^{(t-1)}; \boldsymbol{\theta}) = \sigma(\mathbf{W}_{\text{rec}}\mathbf{h}^{(t-1)} + \mathbf{W}_{\text{in}}\mathbf{x}^{(t)})$$


- \mathbf{W}_{rec} and \mathbf{W}_{in} are parameters
- $\mathbf{h}^{(t)}$ (hidden) states at time t
- σ can be the relu, hyperbolic tangent, or anything


more
trad

Recurrent unit



Recurrent unit




Recurrent unit (with linear activation)

- Recurrent unit


$$\mathbf{h}^{(t)} = f(\mathbf{x}^{(t)}, \mathbf{h}^{(t-1)}; \boldsymbol{\theta}) = \sigma(\mathbf{W}_{\text{rec}}\mathbf{h}^{(t-1)} + \mathbf{W}_{\text{in}}\mathbf{x}^{(t)})$$

```
1 dim_hidden = 6
2 dim_input = 1
3 dim_output = 1
4
5 Wrec = torch.randn(dim_hidden, dim_hidden, requires_grad=True)
6 Win  = torch.randn(dim_hidden, dim_input, requires_grad=True)
7 Wout = torch.randn(dim_hidden, dim_output, requires_grad=True)
8
9 def recurrent_unit(x, h_prev):
10     # YOUR CODE HERE
11     raise NotImplementedError
```

Recurrent unit



Computing the maximum of a sequence



Computing the maximum of a sequence

```
1 minibatch_size = 2**8
2 seq_length = 5
3 for epoch in range(num_epochs):
4     loss = 0
5     for _ in range(minibatch_size):
6         xs = torch.tensor(np.random.randn(seq_length, 1), dtype=torch.
float32)
7         max_xs = torch.max(xs).item()
8         yhat = forward(xs, Wrec, Win, Wout)
9         loss += torch.abs(yhat - max_xs) # try the MSE loss!
10    loss = loss / minibatch_size
11
12    optimizer.zero_grad()
13    loss.backward()
14    optimizer.step()
```

"seq" = $x_s = [6, 2, \dots]$

↓

← $\max(\text{seq})$

← \hat{y}

$(\hat{y} - \max_{x_s}) ** 2$

In class exercise 2

torch.relu

$$h^{(t)} = f(x^{(t)}, h^{(t-1)}; \theta) = \sigma(\underbrace{W_{\text{rec}} h^{(t-1)}}_{\text{given}} + \underbrace{W_{\text{in}} x^{(t)}}_{\text{given}})$$

```
1 def recurrent_unit(x, h_prev): # returns h_next
2     # YOUR CODE HERE
3     raise NotImplementedError
```

Wrec, Win

task 1

```
4
5
6 def forward(seq, Wrec, Win, Wout): # returns yhat
7     # YOUR CODE HERE
8     # hint: use 'for x in seq: ...'
9     raise NotImplementedError
```

initialize h(0)

Vectorize the minibatch?

triple for loop


```
1 minibatch_size = 2**8 256
2 seq_length = 5
3 for epoch in range(num_epochs):
4     loss = 0
5     # CAN WE GET RID OF THE FOLLOWING FOR loop?
6     for _ in range(minibatch_size):
7         xs = torch.tensor(np.random.randn(seq_length, 1), dtype=torch.
float32)
8         max_xs = torch.max(xs).item()
9         yhat = forward(xs, Wrec, Win, Wout)
10        loss += torch.abs(yhat - max_xs)
11    loss = loss / minibatch_size
12
13    optimizer.zero_grad()
14    loss.backward()
15    optimizer.step()
```

vectorized
the
minibatch

In class exercise 3

```
1 def forward(seqs, Wrec, Win, Wout):  
2     """  
3     INPUT  
4     seqs - a (minibatch_size, seq_length, dim_input) tensor  
5  
6     RETURN  
7     yhat - a (minibatch_size, ) tensor  
8     """  
9     # YOUR CODE HERE  
10    raise NotImplementedError
```

Possible to
achieve
zero error
(challenge)



In class exercise 3

Expressivity

↳ universal
↳ approximate

gap

achievable
via
SGD
adam

```
1 def forward(seqs, Wrec, Win, Wout):  
2     """  
3     INPUT  
4     seqs - a (minibatch_size, seq_length, dim_input) tensor  
5  
6     RETURN  
7     yhat - a (minibatch_size, ) tensor  
8     """  
9     # YOUR CODE HERE  
10    raise NotImplementedError
```

there exists a set of weight (W_{rec}, W_{in}, W_{out})
st. we can get 0 error (forever)
Upshot: RNNs are very expressive "universal approximation"

An annoying loop

```
1 def forward(seqs, Wrec, Win, Wout):
2     """
3     INPUT
4     seqs - a (minibatch_size, seq_length, dim_input) tensor
5
6     RETURN
7     yhat - a (minibatch_size, ) tensor
8     """
9     h = torch.zeros(seqs.shape[0], dim_hidden)
10
11     for x in seqs.transpose(0, 1): # CAN WE GET RID OF THIS?
12         h = recurrent_unit(x, h, Wrec, Win)
13     # [...]
```

Vectorize
over time?

- No, at least not yet (active research area)

An annoying loop

```
1 def forward(seqs, Wrec, Win, Wout):
2     """
3     INPUT
4     seqs - a (minibatch_size, seq_length, dim_input) tensor
5
6     RETURN
7     yhat - a (minibatch_size, ) tensor
8     """
9     h = torch.zeros(seqs.shape[0], dim_hidden)
10
11     for x in seqs.transpose(0, 1): # CAN WE GET RID OF THIS?
12         h = recurrent_unit(x, h, Wrec, Win)
13     # [...]
```

"(# of param)
square"

not SGD
not adam
need 2nd deriv
Newton's

Catch

- No, at least not yet (active research area)
- "Parallelizing non-linear sequential models over the sequence length" (Lim et al., 2024 ICLR)

$\sigma \neq \text{identity}$

An annoying loop


```
1 def forward(seqs, Wrec, Win, Wout):  
2     """  
3     INPUT  
4     seqs - a (minibatch_size, seq_length, dim_input) tensor  
5  
6     RETURN  
7     yhat - a (minibatch_size, ) tensor  
8     """  
9     h = torch.zeros(seqs.shape[0], dim_hidden)  
10  
11     for x in seqs.transpose(0, 1): # CAN WE GET RID OF THIS?  
12         h = recurrent_unit(x, h, Wrec, Win)  
13     # [...]
```

$\frac{2f}{\partial \theta_i}$ first order
 $i \in \{\text{all params}\}$
 $(\# \text{ param})^2 = \frac{2^2 f}{\partial \theta_i \partial \theta_j}$
 $i, j \in \{\text{all params}\}$

- No, at least not yet (active research area)
- “Parallelizing non-linear sequential models over the sequence length” (Lim et al., 2024 ICLR)

↳ Neural ODE apply DL to physics

Recurrent Neural Network with multiple layers



Implementation

```
1 class recurrent_cell(nn.Module):
2     def __init__(self, input_dim, hidden_dim):
3         # [...]
4         self.hidden_dim = hidden_dim
5         self.input_to_hidden = nn.Linear(input_dim, hidden_dim)
6         self.hidden_to_hidden = nn.Linear(hidden_dim, hidden_dim)
7
8     def forward(self, x, h_prev):
9         h_next = torch.relu(self.input_to_hidden(x) + self.
hidden_to_hidden(h_prev))
10        return h_next
```

Implementation

```
1 class RNN_layer(nn.Module):
2     def __init__(self, embed_dim, hidden_dim):
3         # [...]
4         self.rnn_layer = recurrent_cell(embed_dim, hidden_dim)
5
6     def forward(self, x_seq, h):
7         outputs = []
8         for t in range(x_seq.size(1)):
9             h = self.rnn_layer(x_seq[:, t, :], h)
10            outputs.append(h)
11
12        # Stack outputs
13        x_transformed = torch.stack(outputs, dim=1)
14        return x_transformed
```


Implementation

```
1 class RNN(nn.Module):
2     def __init__(self, vocab_size, embed_dim, hidden_dim, num_layers):
3         # [...]
4         self.embedding = nn.Embedding(vocab_size, embed_dim)
5         self.layers = nn.ModuleList([RNN_layer(embed_dim, hidden_dim)
6 for _ in range(num_layers)])
7         self.classification_head = nn.Linear(embed_dim, vocab_size)
8
9     def forward(self, x):
10         x = self.embedding(x)
11         h = torch.zeros(x.size(0), self.layers[0].rnn_layer.hidden_dim,
12 device=x.device) # Initial hidden state
13
14         for rnn_layer in self.layers:
15             x = rnn_layer(x, h)
16
17         logits = self.classification_head(x)
18         return logits
```


Seq-2-“one” Causal models

$$f(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(t)}; \theta) = \mathbf{y}^{(t)}$$

Causality assumption

The output $\mathbf{y}^{(t)}$ does not depend on future values of $\mathbf{x}^{(t')}$ where $t' > t$

Seq-2-seq causal models

$$h^{(1)} \dots h^{(C)}$$

$$f(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(C)}; \theta) = \mathbf{y}^{(1)}, \dots, \mathbf{y}^{(C)}$$

translation

Causality assumption

The output $\mathbf{y}^{(t)}$ does not depend on future values of $\mathbf{x}^{(t')}$ where $C \geq t' > t \geq 1$

transformer \rightarrow self-attention is a Seq2seq
multilayer model
Does not suffer from non-parallelization in time/seq length

Sequences

linear algebra-ification

$$\mathbf{X} = [\mathbf{x}^{(1)} \quad \mathbf{x}^{(2)} \quad \dots \quad \mathbf{x}^{(C)}] \in \mathbb{R}^{d \times C}$$

where

- $\mathbf{x}^{(t)} \in \mathbb{R}^d$
- $t \in \{1, \dots, C\}$,
- we should really write $\mathbf{X}^{(i)}$ where $i \in \{1, \dots, N\}$ but hide from notation for convenience

$\mathbf{X}^{(i)}$

Self-attention

$X \neq Z$ self $X = Y = Z$
can't plug
in different
things

- $\text{self-attention}(\mathbf{X}; \theta) := \underbrace{\mathbf{W}^{(V)\top}}_{(d, d)} \mathbf{X} \underbrace{\text{softmax} \left(\underbrace{\mathbf{X}^\top \mathbf{W}^{(K)\top}}_{(d, d)} \underbrace{\mathbf{W}^{(Q)}}_{(d, d)} \mathbf{X} \right)}_{(d, d)} \in \mathbb{R}^{d \times C}.$

- parameters

$$\theta^{(\text{att})} = [\mathbf{W}^{(Q)}, \mathbf{W}^{(K)}, \mathbf{W}^{(V)}]$$

- (“Q, K, V” stands for “query”, “key”, “value”, respectively)
where

$$\mathbf{W}^{(Q)} \text{ and } \mathbf{W}^{(K)} \text{ and } \mathbf{W}^{(V)} \in \mathbb{R}^{d \times d}.$$

- “seq-2-seq”: maps the sequence \mathbf{X} to another sequence $\text{attention}(\mathbf{X}; \theta)$:
- **Problem:** is this causal?

$\underbrace{\text{softmax}(K^\top Q)}_{(C, C)}$ - column stochastic
col sum to 1

Self-attention

$$\begin{matrix} (d, d) & (d, C) & [C, C] \\ \underbrace{\hspace{1cm}} & \underbrace{\hspace{1cm}} & \underbrace{\hspace{2cm}} \end{matrix}$$

- $\text{attention}(\mathbf{X}; \theta) := \mathbf{W}^{(V)\top} \mathbf{X}^{(i)} \text{softmax} \left(\mathbf{X}^\top \mathbf{W}^{(K)\top} \mathbf{W}^{(Q)} \mathbf{X} \right) \in \mathbb{R}^{d \times C}.$

- parameters

$$\theta^{(\text{att})} = [\mathbf{W}^{(Q)}, \mathbf{W}^{(K)}, \mathbf{W}^{(V)}]$$

- (“Q, K, V” stands for “query”, “key”, “value”, respectively)
where

$$\mathbf{W}^{(Q)} \text{ and } \mathbf{W}^{(K)} \text{ and } \mathbf{W}^{(V)} \in \mathbb{R}^{d \times d}.$$

- “seq-2-seq”: maps the sequence \mathbf{X} to another sequence $\text{attention}(\mathbf{X}; \theta)$:

- Problem:** is this causal?

$$\text{input.shape} = \text{output.shape}$$

Is this causal?

$$\exp KQ$$

$$\text{softmax}(K^T Q)$$


$C=4$

$$\text{attention}(X) =$$

$$W_V \begin{bmatrix} x_0 & x_1 & x_2 & x_3 \end{bmatrix}$$

X

start here
my rec



In class exercise 4

$$\begin{pmatrix} \text{ } & \text{ } & \text{ } & \text{ } \end{pmatrix}$$

$\text{att}(X)[0,2]$

$=$

$$W_V (p_0 x_0 + p_1 x_1 + p_2 x_2 + p_3 x_3)$$

$$\begin{bmatrix} p_0 \\ p_1 \\ p_2 \\ p_3 \end{bmatrix}$$

$\neq 0$

Is this causal?

$$W_V \times \left(\text{softmax}(KQ) * M \right)$$


In class exercise 4

not quite
right

but on
right
track

Causal mask

Flash attention



**Memory Hierarchy with
Bandwidth & Memory Size**

From Dao et al 2022 (Flash Attention)

Flash attention

Algorithm 0 Standard Attention Implementation

Require: Matrices $\mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{N \times d}$ in HBM.

- 1: Load \mathbf{Q}, \mathbf{K} by blocks from HBM, compute $\mathbf{S} = \mathbf{QK}^\top$, write \mathbf{S} to HBM.
 - 2: Read \mathbf{S} from HBM, compute $\mathbf{P} = \text{softmax}(\mathbf{S})$, write \mathbf{P} to HBM.
 - 3: Load \mathbf{P} and \mathbf{V} by blocks from HBM, compute $\mathbf{O} = \mathbf{PV}$, write \mathbf{O} to HBM.
 - 4: Return \mathbf{O} .
-

From Dao et al 2022 (Flash Attention)

References I
